

Elasticsearch In Netflix

Danny Yuan, Jae Bae



Welcome



Hashtag: #ES_in_Netflix

@Elasticsearch - Elasticsearch

A screenshot of the Netflix Open Source GitHub repository page. The header includes the Netflix OSS logo and links for Repositories, Commit Timeline, Mailing Lists, Community, and Powered By. Below the header, there's a section titled "Our Repositories" with three cards: "Archaius" (ANIME & ANIMATION), "Asgard" (FAMILY ANIMATION), and "Astyanax" (SCI-FI & FANTASY). Each card has a thumbnail image and the repository name.

@stonse - Sudhir Tonse

@g9yuayon - Danny Yuan

@metacret - Jae Bae

Who Are We?



Who Are We?

Software engineers in Netflix's Platform Engineering team, working on large scale data infrastructure



Who Are We?

Software engineers in Netflix's Platform Engineering team, working on large scale data infrastructure

Building and operating Netflix's cloud real-time query service



Why Are We Here?



Why Are We Here?

How We Use Elasticsearch



Why Are We Here?

How We Use Elasticsearch

Why Elasticsearch



Why Are We Here?

How We Use Elasticsearch

Why Elasticsearch

How We Run Elasticsearch



Why Are We Here?

How We Use Elasticsearch

Why Elasticsearch

How We Run Elasticsearch

To Seek Your Feedback



How We Use Elasticsearch



Querying Log Events

Tracking Service Deployments

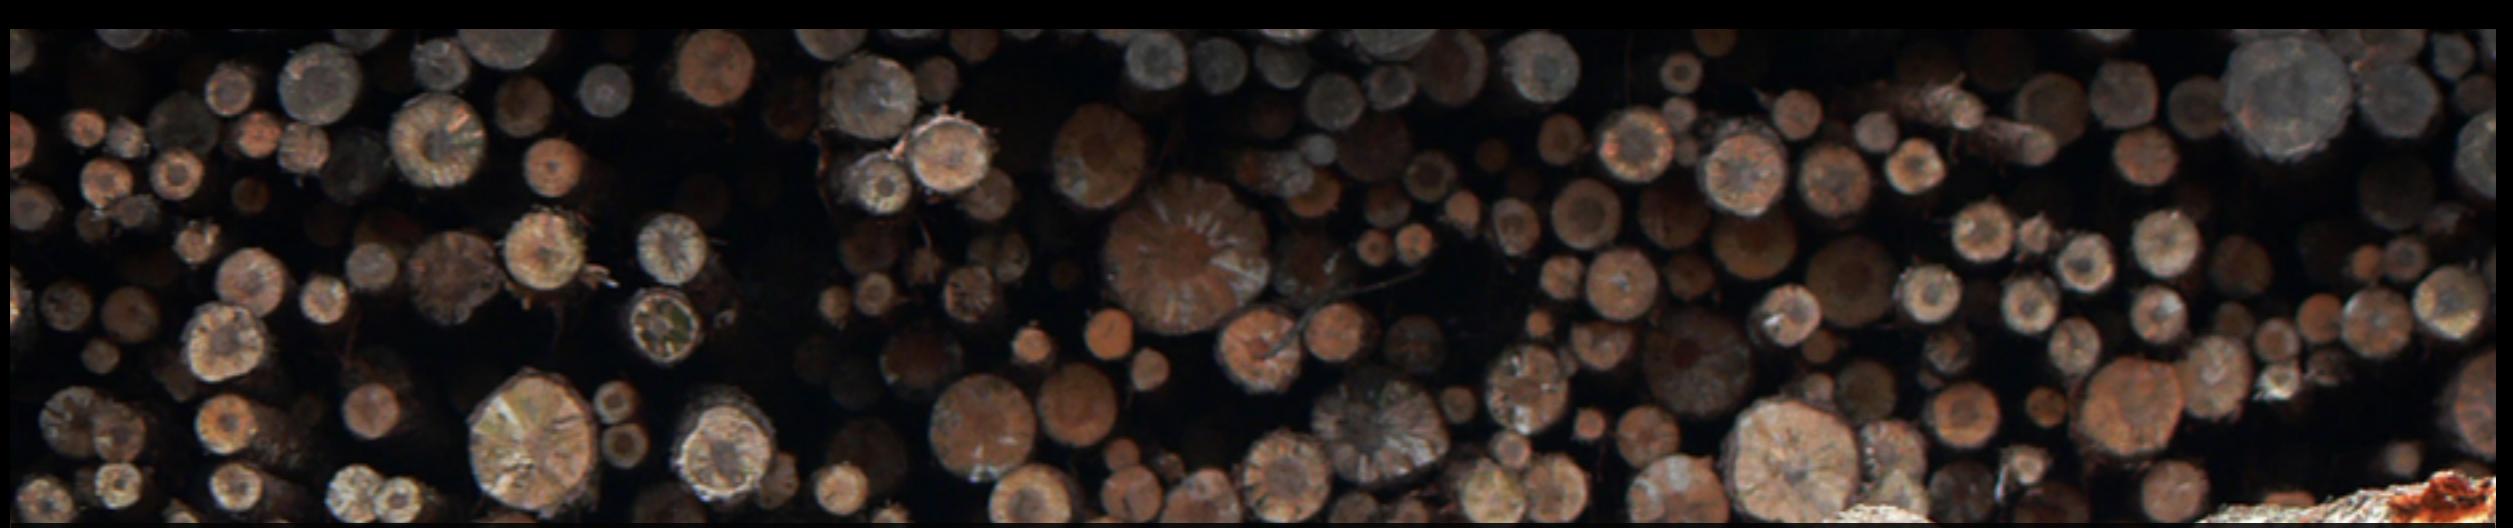


Querying Log Events



A Little Historical Perspective





**Netflix is a log generating company
that also happens to stream movies**

- Adrian Cockcroft



photo credit: http://www.flickr.com/photos/decade_null/142235888/sizes/o/in/photostream/





NETFLIX

Operational Insights Business Analysis Debugging Trend



A Humble Beginning



A Humble Beginning

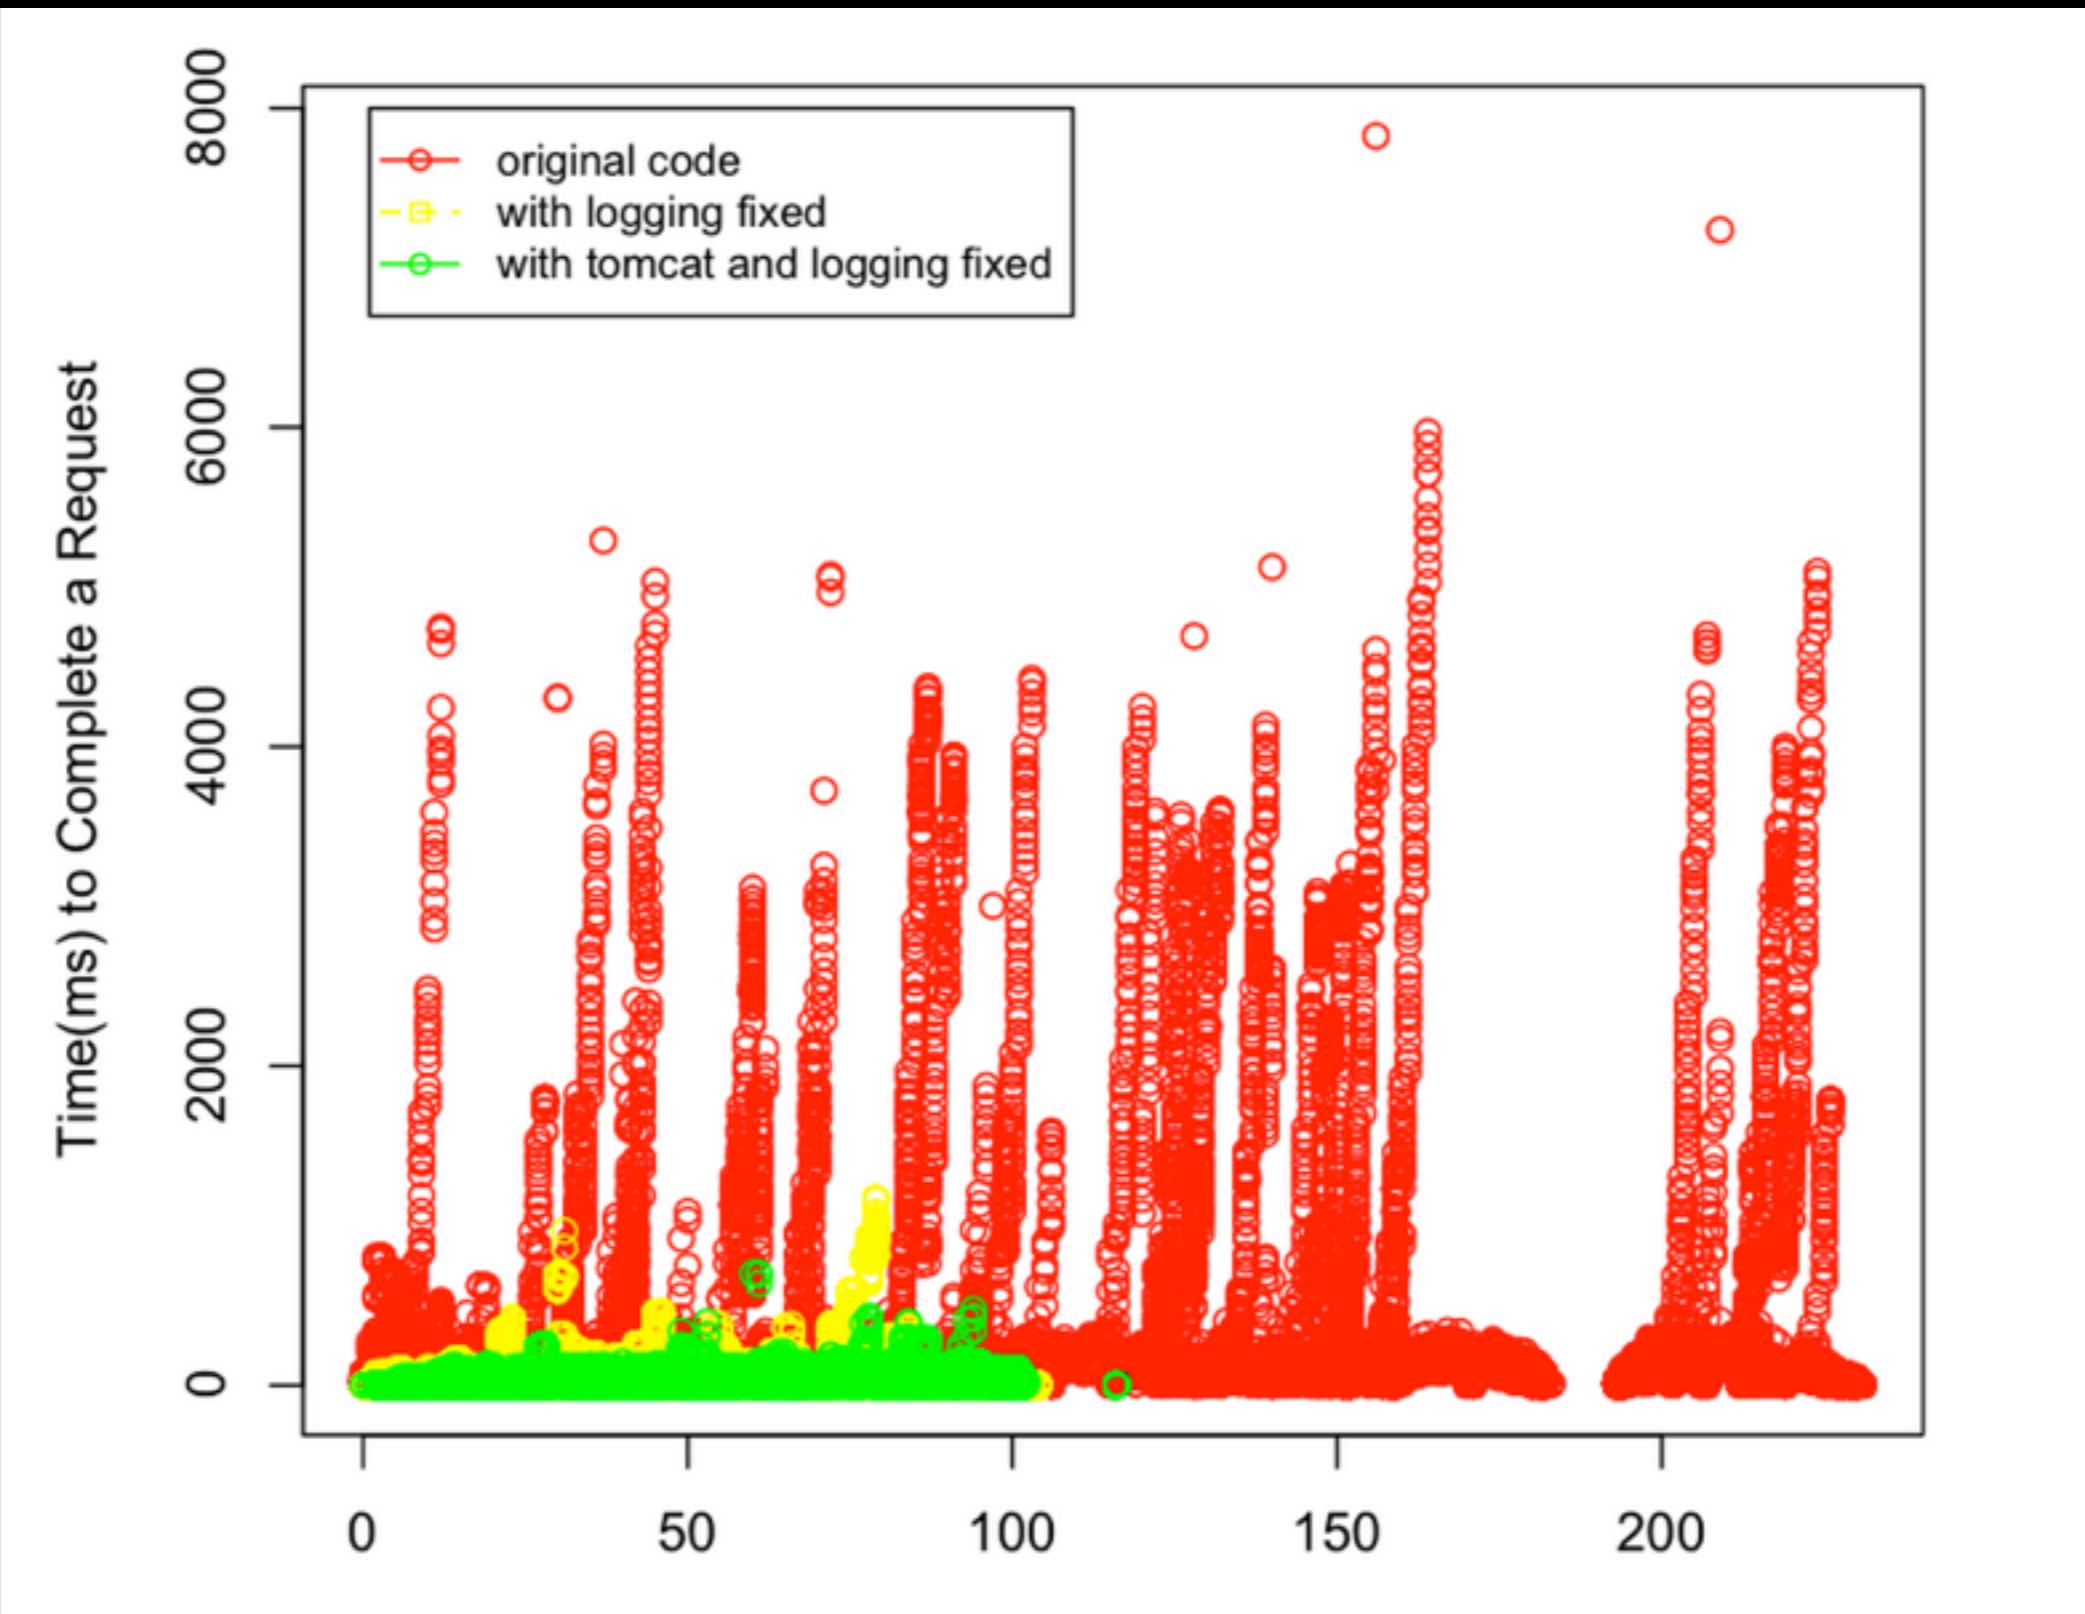
```
awk '{print $3, $10}' access.log | sort | uniq -c | sort -nr
```



A Humble Beginning

```
plot_all <- function(files, colors){  
  file <- head(files, 1)  
  color <- head(colors, 1)  
  data <- load_data(file)  
  plot(data, type="p", col=color, xlab="Elapsed Time (ms)", ylab="Ti  
max_latency = max(data$V2)  
z = ?l
```



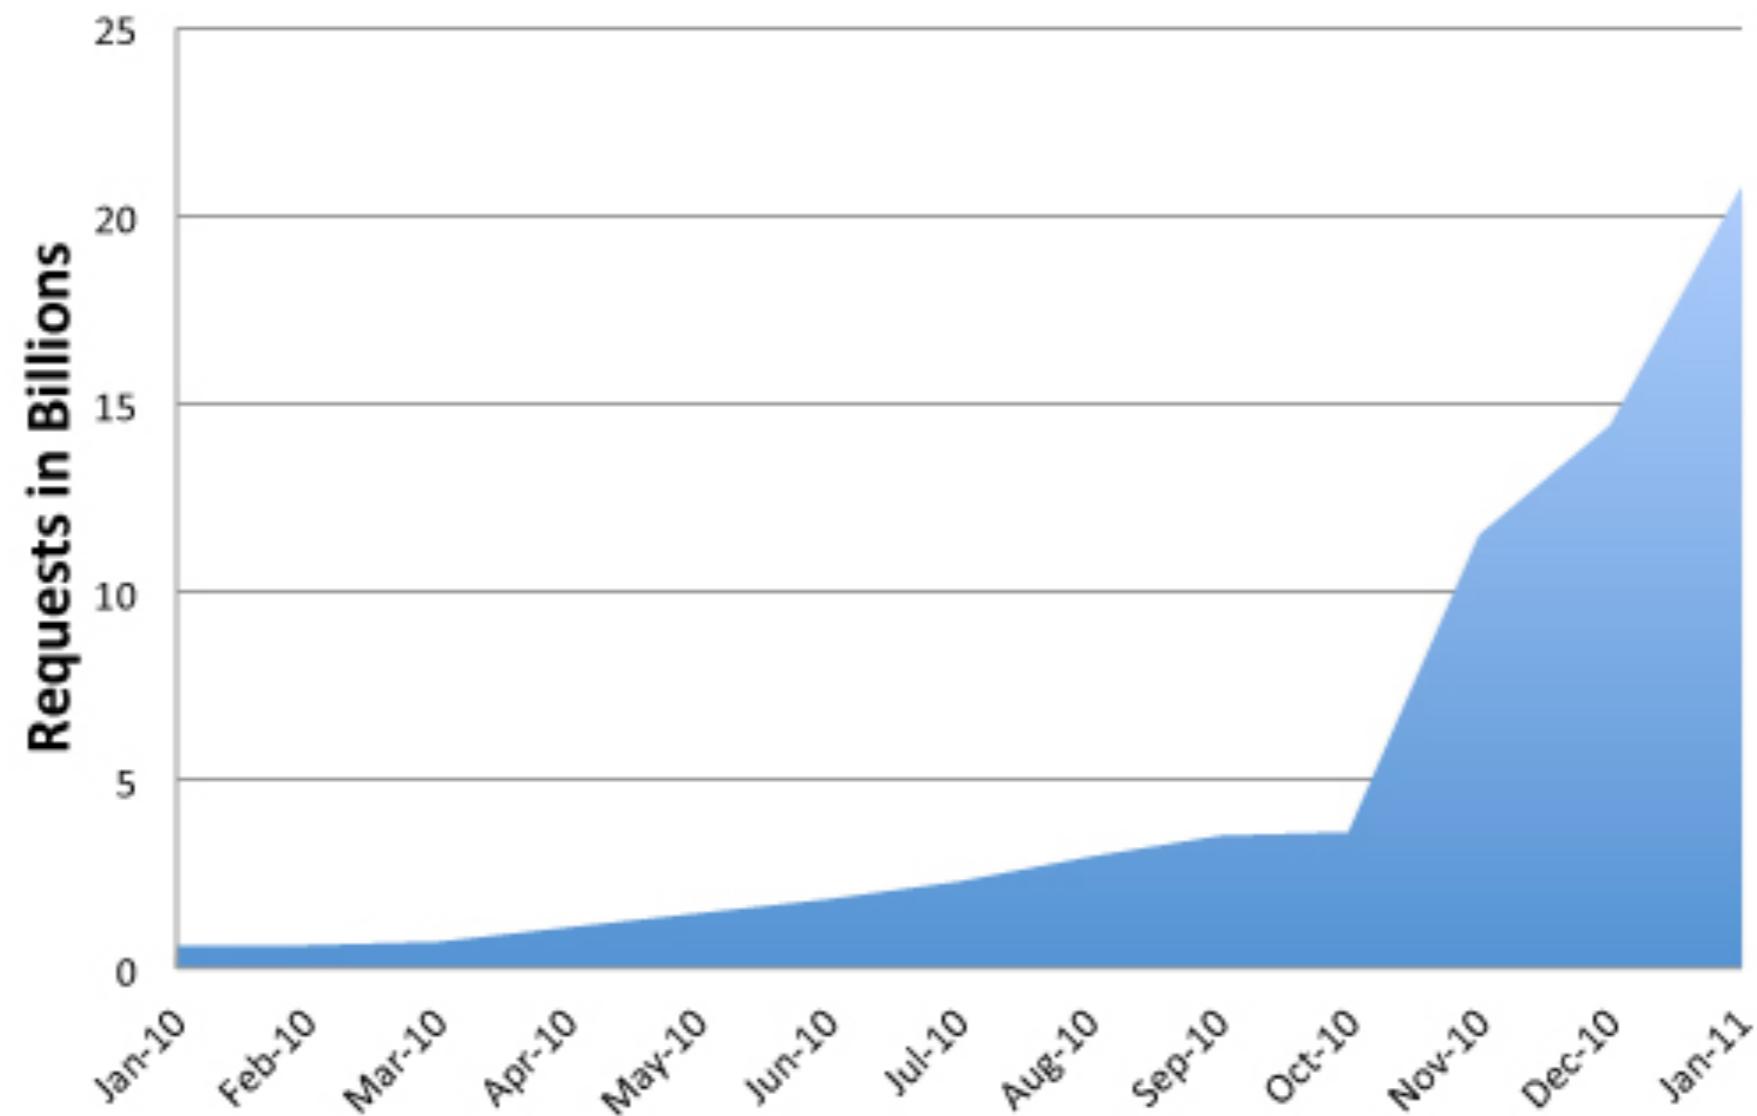


Things Changed



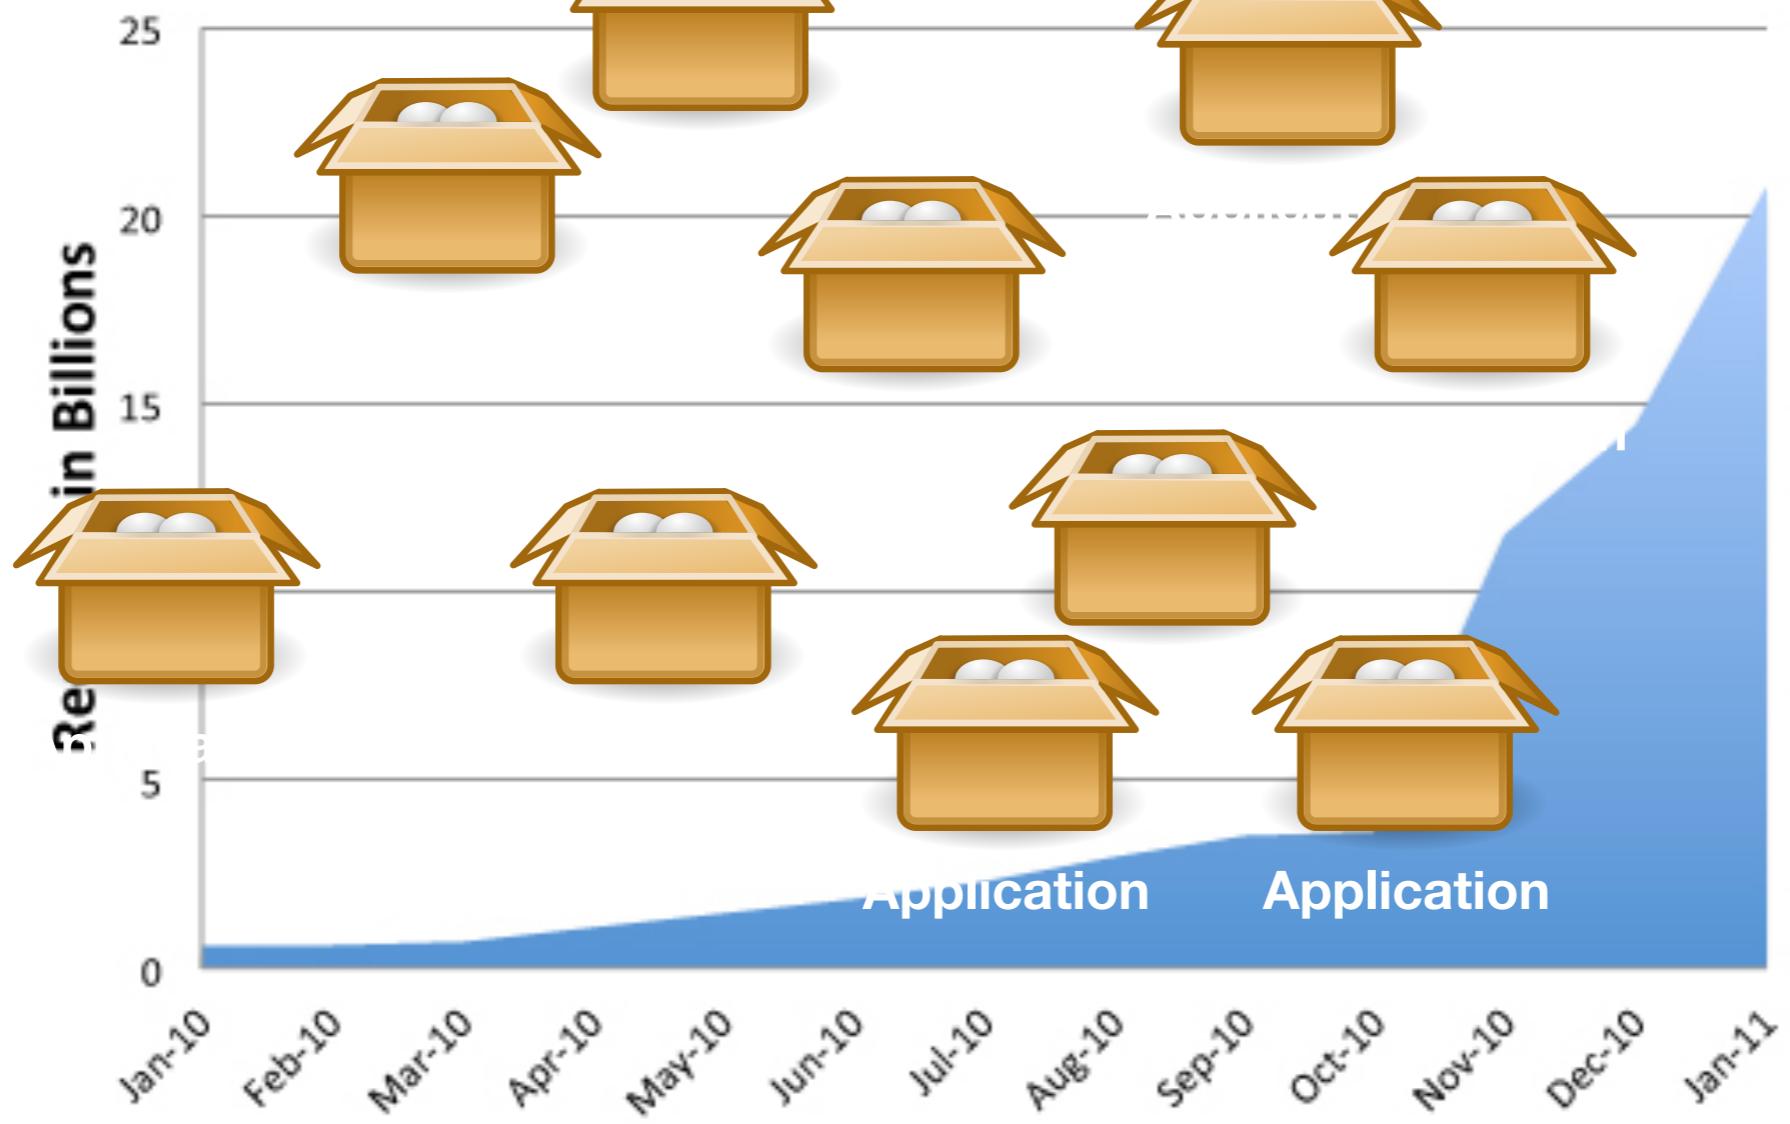


Netflix API : Growth in Requests



NETFLIX®

Netflix API Growth in Requests



NETFLIX®

70,000,000,000



1,500,000



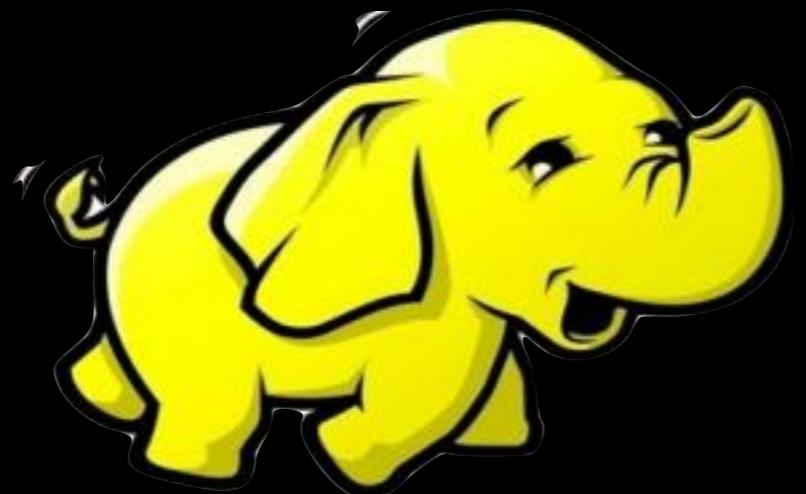
Making Sense of Billions of Events

So We Evolved



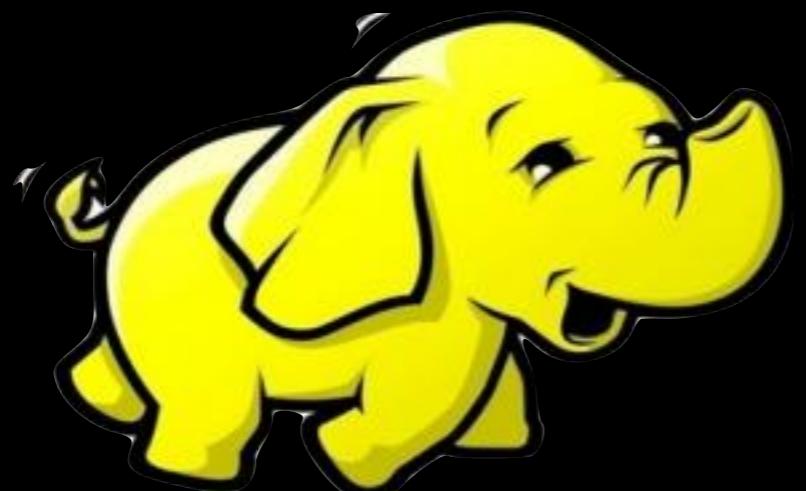
So We Evolved

hadoop



So We Evolved

hadoop

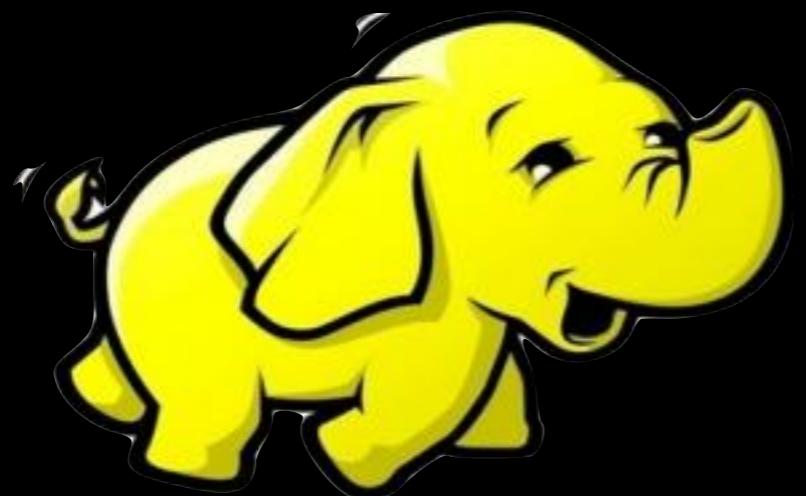


```
hgrep -C 10 -k 5,2,3 'users.*[l-9]{3}' *catalina.out s3//bucket
```



So We Evolved

hadoop

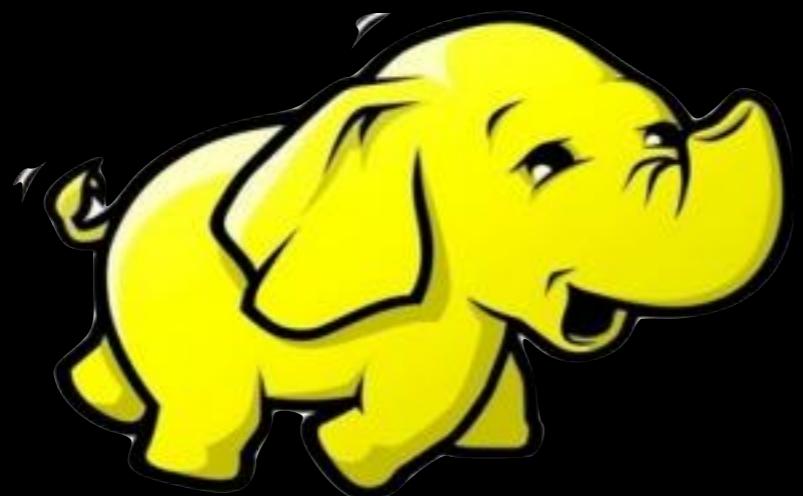


```
hgrep -C 10 -k 5,2,3 'users.*[1-9]{3}' *catalina.out s3//bucket
```



So We Evolved

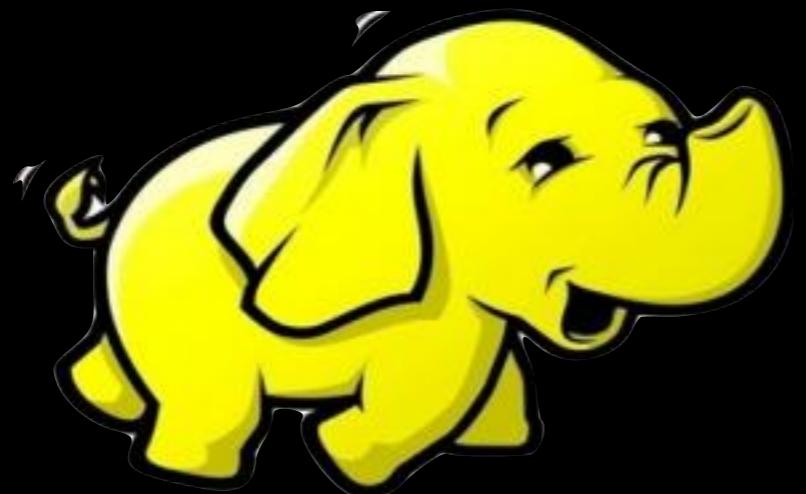
hadoop



```
hgrep -C 10 -k 5,2,3 'users.*[!-9]{3}' *catalina.out s3//bucket
```

So We Evolved

hadoop

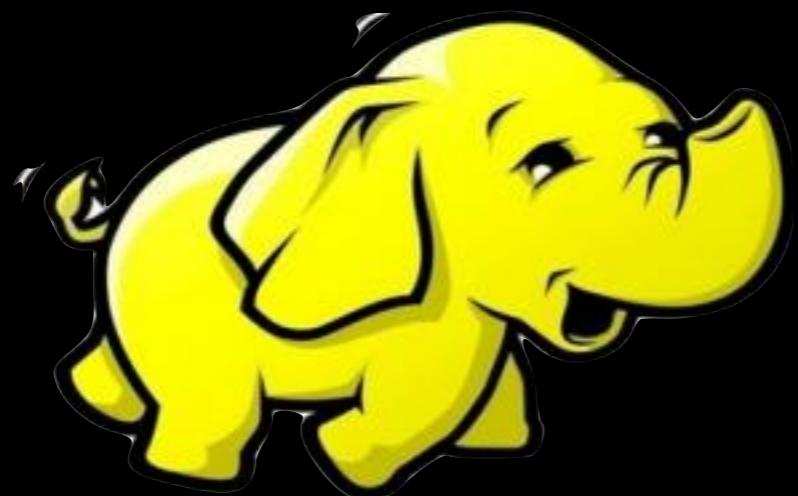


```
hgrep -C 10 -k 5,2,3 'users.*[!-9]{3}' *catalina.out s3//bucket
```



So We Evolved

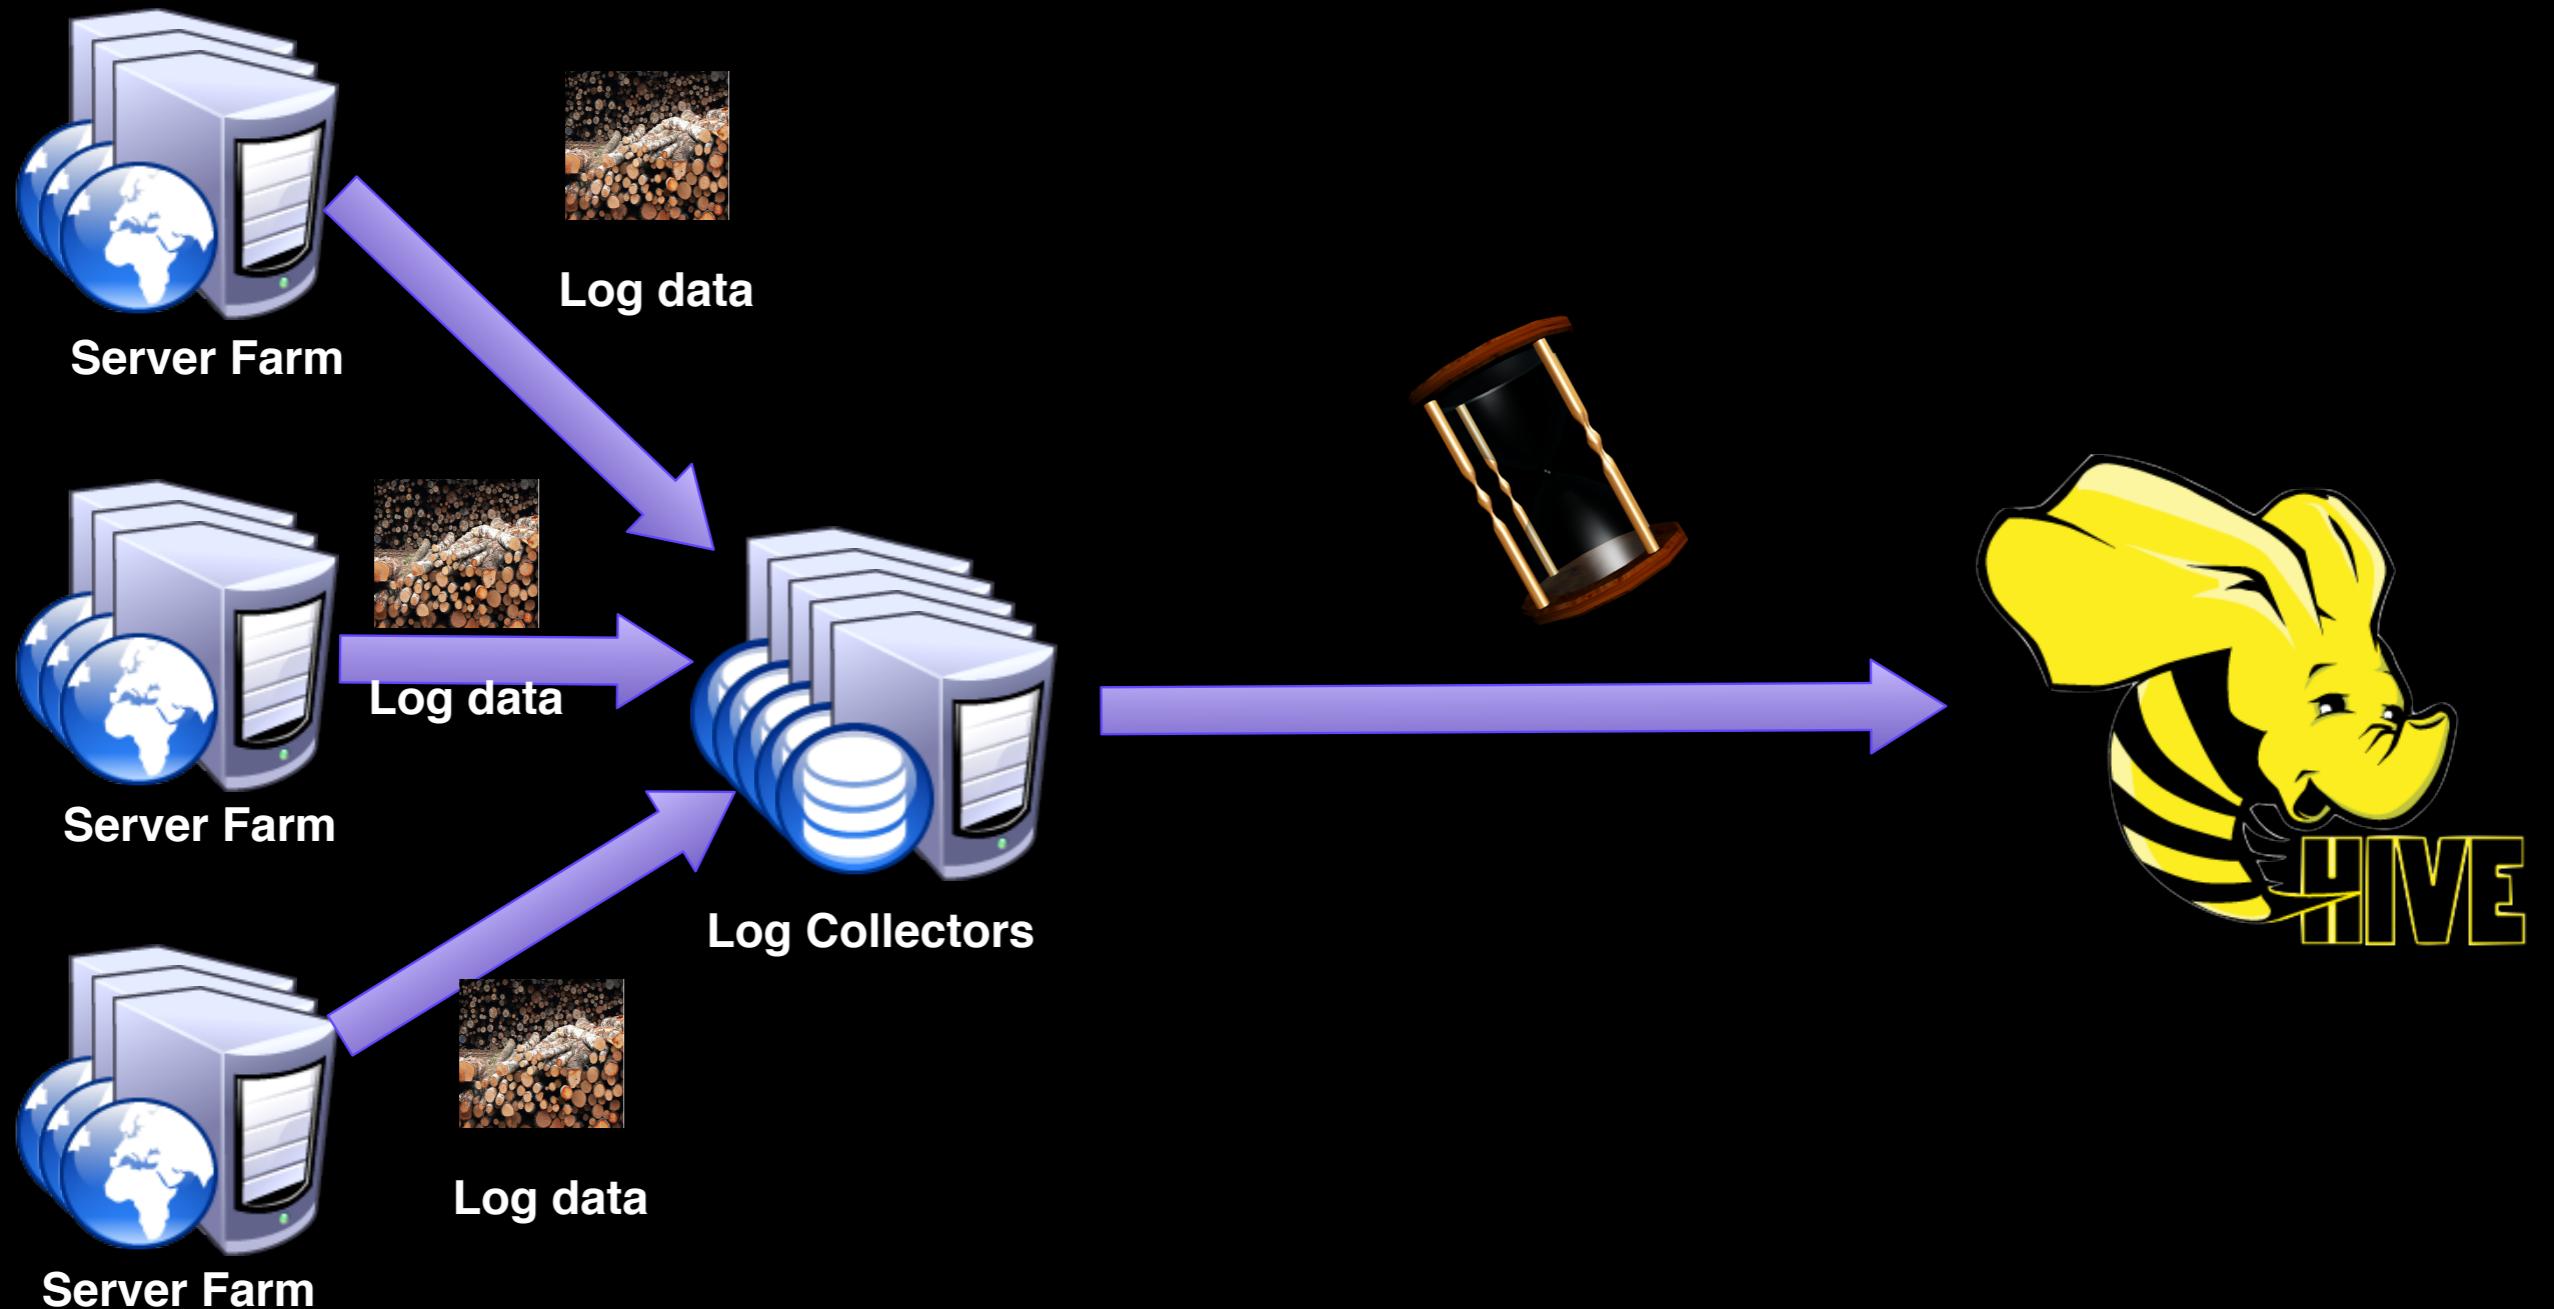
hadoop



```
hgrep -C 10 -k 5,2,3 'users.*[l-9]{3}' *catalina.out s3//bucket  
select * from log_events where dateint=20140101
```

Field Name	Field Value
Client	“API”
Server	“Cryptex”
StatusCode	200
ResponseTime	73





NETFLIX®

I have an issue!

Biosys Document

Home

Biosys Test US-EAST-1 ▾

Instance Information

ASG: cryptex_api-v006						Deselect All Instances
Show	10	entries	Search: b252b5c1			
ASG Name	Instance ID	Type	Zone	Start	State	
<input checked="" type="checkbox"/> cryptex_api-v006	i-dcc845af	m1.large	us-east-1e	2013-02-25 21:57:21.503 UTC	InService	
<input checked="" type="checkbox"/> cryptex_api-v006	i-2c4d975f	m1.large	us-east-1d	2013-02-21 21:39:43.185 UTC	InService	

Showing 1 to 2 of 2 entries

First Previous 1 Next Last

Search Options

File Pattern: catalina.out
Pattern: Exception

Show File Name: Only File Name: Show Line No:
Context: Before: After:
Max Lines: Ignore Case? Any Time [?](#)
Past Hours: (PST-0800): 02/26 21:03 – 02/25 21:03 (24.0 hours)

Download Search Result [Search](#) [Show Files](#)

File List

Search Result

Tail

Search Status

[Cancel Search](#) 37 lines are found Streaming is done.

Search Result

```
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: java.net.SocketTimeoutException: Read timed out
i-2c4d975f: Caused by: java.net.SocketTimeoutException: Read timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: java.net.SocketTimeoutException: Read timed out
i-2c4d975f: Caused by: java.net.SocketTimeoutException: Read timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: java.net.SocketTimeoutException: Read timed out
i-2c4d975f: Caused by: java.net.SocketTimeoutException: Read timed out
i-2c4d975f: java.lang.NullPointerException
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-dec845af: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-50-19-255-91.compute-1.amazonaws.com:7001 timed out
i-dec845af: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-50-19-255-91.compute-1.amazonaws.com:7001 timed out
```



I have an issue!

Biosys Document

Home

Biosys Test US-EAST-1 ▾

Instance Information

ASG: cryptex_api-v006			
Show 10 entries			
▲	ASG Name	Instance ID	Type
<input checked="" type="checkbox"/>	cryptex_api-v006	i-dcc845af	m1.large
<input checked="" type="checkbox"/>	cryptex_api-v006	i-2c4d975f	m1.large

Showing 1 to 2 of 2 entries

INSTANCE INFORMATION

ASG: cryptex_api-v006						Deselect All Instances
Show 10 entries						Search: b252b5c1
▲	ASG Name	Instance ID	Type	Zone	Start	State
<input checked="" type="checkbox"/>	cryptex_api-v006	i-dcc845af	m1.large	us-east-1e	2013-02-25 21:57:21.503 UTC	InService
<input checked="" type="checkbox"/>	cryptex_api-v006	i-2c4d975f	m1.large	us-east-1d	2013-02-21 21:39:43.185 UTC	InService

Showing 1 to 2 of 2 entries

First Previous 1 Next Last

File List

Search Result

Tail

Search Status

[Cancel Search](#) 37 lines are found Streaming is disabled

Search Result

i-2c4d975f: com.sun.jersey.api.client
i-2c4d975f: Caused by: java.net.SocketException
i-2c4d975f: com.sun.jersey.api.client
i-2c4d975f: Caused by: java.net.SocketException
i-2c4d975f: com.sun.jersey.api.client
i-2c4d975f: Caused by: java.net.SocketException
i-2c4d975f: java.lang.NullPointerException
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-dec845af: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-50-19-255-91.compute-1.amazonaws.com:7001 timed out
i-dec845af: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-50-19-255-91.compute-1.amazonaws.com:7001 timed out



Instance Information

ASG: cryptex_api-v006 ✘

Show 10 : entries

	ASG Name	Instance ID	Type
<input checked="" type="checkbox"/>	cryptex_api-v006	i-dcc845af	m1.large
<input checked="" type="checkbox"/>	cryptex_api-v006	i-2c4d975f	m1.large

Showing 1 to 2 of 2 entries

[File List](#)

Search Result

Search Status

[Cancel Search](#) 37 lines are found Streaming is d

Search Options

File Pattern: catalina.out

Pattern: **Exception**

Show File Name: Only File Name: Show Line No:

Context: 0 Before: 0 After: 0

Max Lines:  Ignore Case? Any Time (?)

Past Hours:

(PST-0800): 02/26 21:07 – 02/25 21:07 (24.0 hours)

(PST-0800): 02/26 21:07 - 02/25 21:07 (24.0 hours)



I have an issue!

Biopsys Document

Home

Biopsys Test US-EAST-1 ▾

Instance Information

ASG: cryptex_ap

Show 10 ▾

▲ ASG Nam

cryptex_a

v006

cryptex_a

v006

Showing 1 to 2 of

Search Options

Search Status

[Cancel Search](#)

37 lines are found Streaming is done.

Search Result

i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException:
i-2c4d975f: Caused by: java.net.SocketTimeoutException: Read
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException:
i-2c4d975f: Caused by: java.net.SocketTimeoutException: Read
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException:
i-2c4d975f: Caused by: java.net.SocketTimeoutException: Read
i-2c4d975f: java.lang.NullPointerException
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException:
170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException:
Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-50-19-255-91.compute-1.amazonaws.com:7001 timed out
Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-50-19-255-91.compute-1.amazonaws.com:7001 timed out



I have an issue!

Biopsys Test US-EAST-1

Biopsys

Home

Tail Controls

Ins

ASQ

Show

Sh

Sh

File List

Sea

Ca

Sea

i-2

Cancel Tail

Select Log Type: NCCP Apache Log

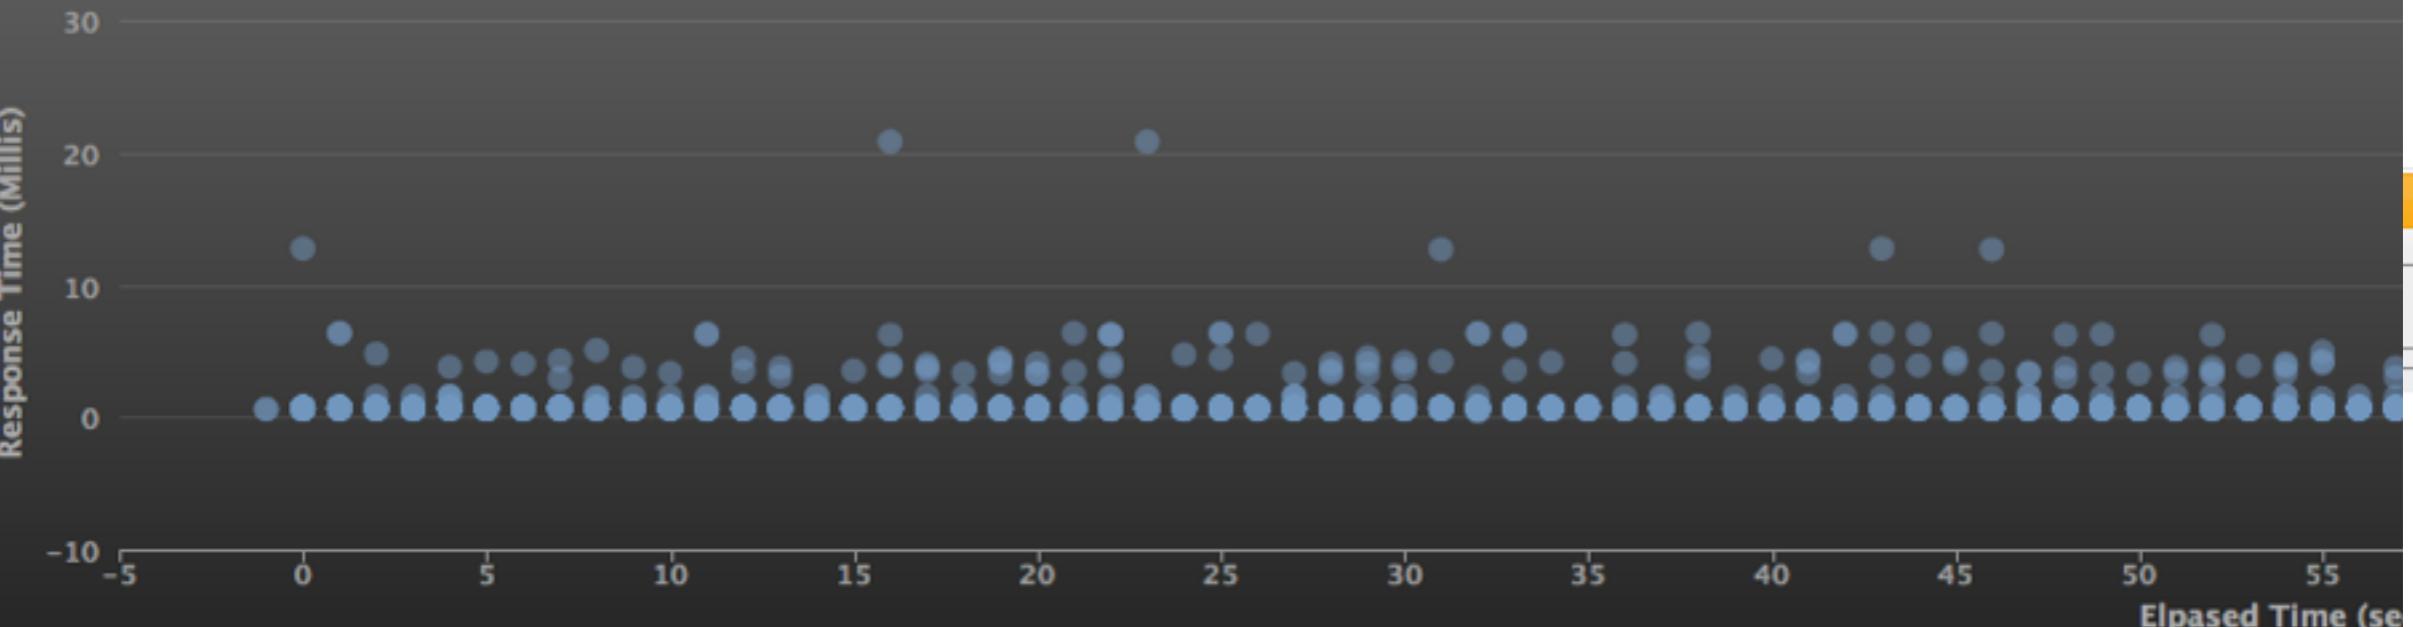
Tail View

Aggregate:

[Clear Results](#)

Access Graph

Response Time versus Elapsed Time



- aggregated
- ◆ /nccp/controller/2.11/apptruststore
- /nccp/controller/2.11/heartbeat
- ▲ /nccp/controller/2.15/heartbeat
- ▼ /nccp/controller/2.12/heartbeat
- ▲ /nccp/controller/2.11/logblob
- ▼ /nccp/controller/2.13/heartbeat
- /nccp/controller/2.15/authenticationrenewal
- ◆ /nccp/controller/2.13/ping
- /nccp/controller/2.12/authenticationrenewal
- /nccp/controller/2.15/playdata
- ▲ /nccp/controller/2.12/logblob
- ▼ /nccp/controller/2.13/authenticationrenewal
- ▲ /nccp/controller/2.12/ping
- ▼ /nccp/controller/2.13/playdata
- /nccp/controller/2.11/authenticationrenewal
- ◆ /nccp/controller/2.15/authorization
- ◆ /nccp/controller/2.15/license
- /nccp/controller/2.11/ping
- ▲ /nccp/controller/2.12/nasverify
- ▼ /nccp/controller/2.13/authorization
- /nccp/controller/2.13/register
- ▲ /nccp/controller/2.14/heartbeat
- ▼ 23.23.50.52
- 23.20.164.0
- ◆ 50.17.70.174
- 50.19.53.131
- ▲ 50.19.9.178

```
Caused by: org.apache.http.ConnectionTimedOutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-2c4d975f: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-204-236-228-170.compute-1.amazonaws.com:7001 timed out
i-dcc845af: com.sun.jersey.api.client.ClientHandlerException: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-50-19-255-91.compute-1.amazonaws.com:7001 timed out
i-dcc845af: Caused by: org.apache.http.conn.ConnectTimeoutException: Connect to ec2-50-19-255-91.compute-1.amazonaws.com:7001 timed out
```



What Could Go Wrong?





"ROUTE FORCE"

NETFLIX®



You thought parallelization would save the day?
Think again

NETFLIX®



$$T(N) = \frac{N}{1 - (1 - B)N}$$

NETFLIX®

What Is Missing?



Interactive Exploration



NETFLIX®

Functional Requirements

Arbitrary Boolean Queries

Aggregated Query

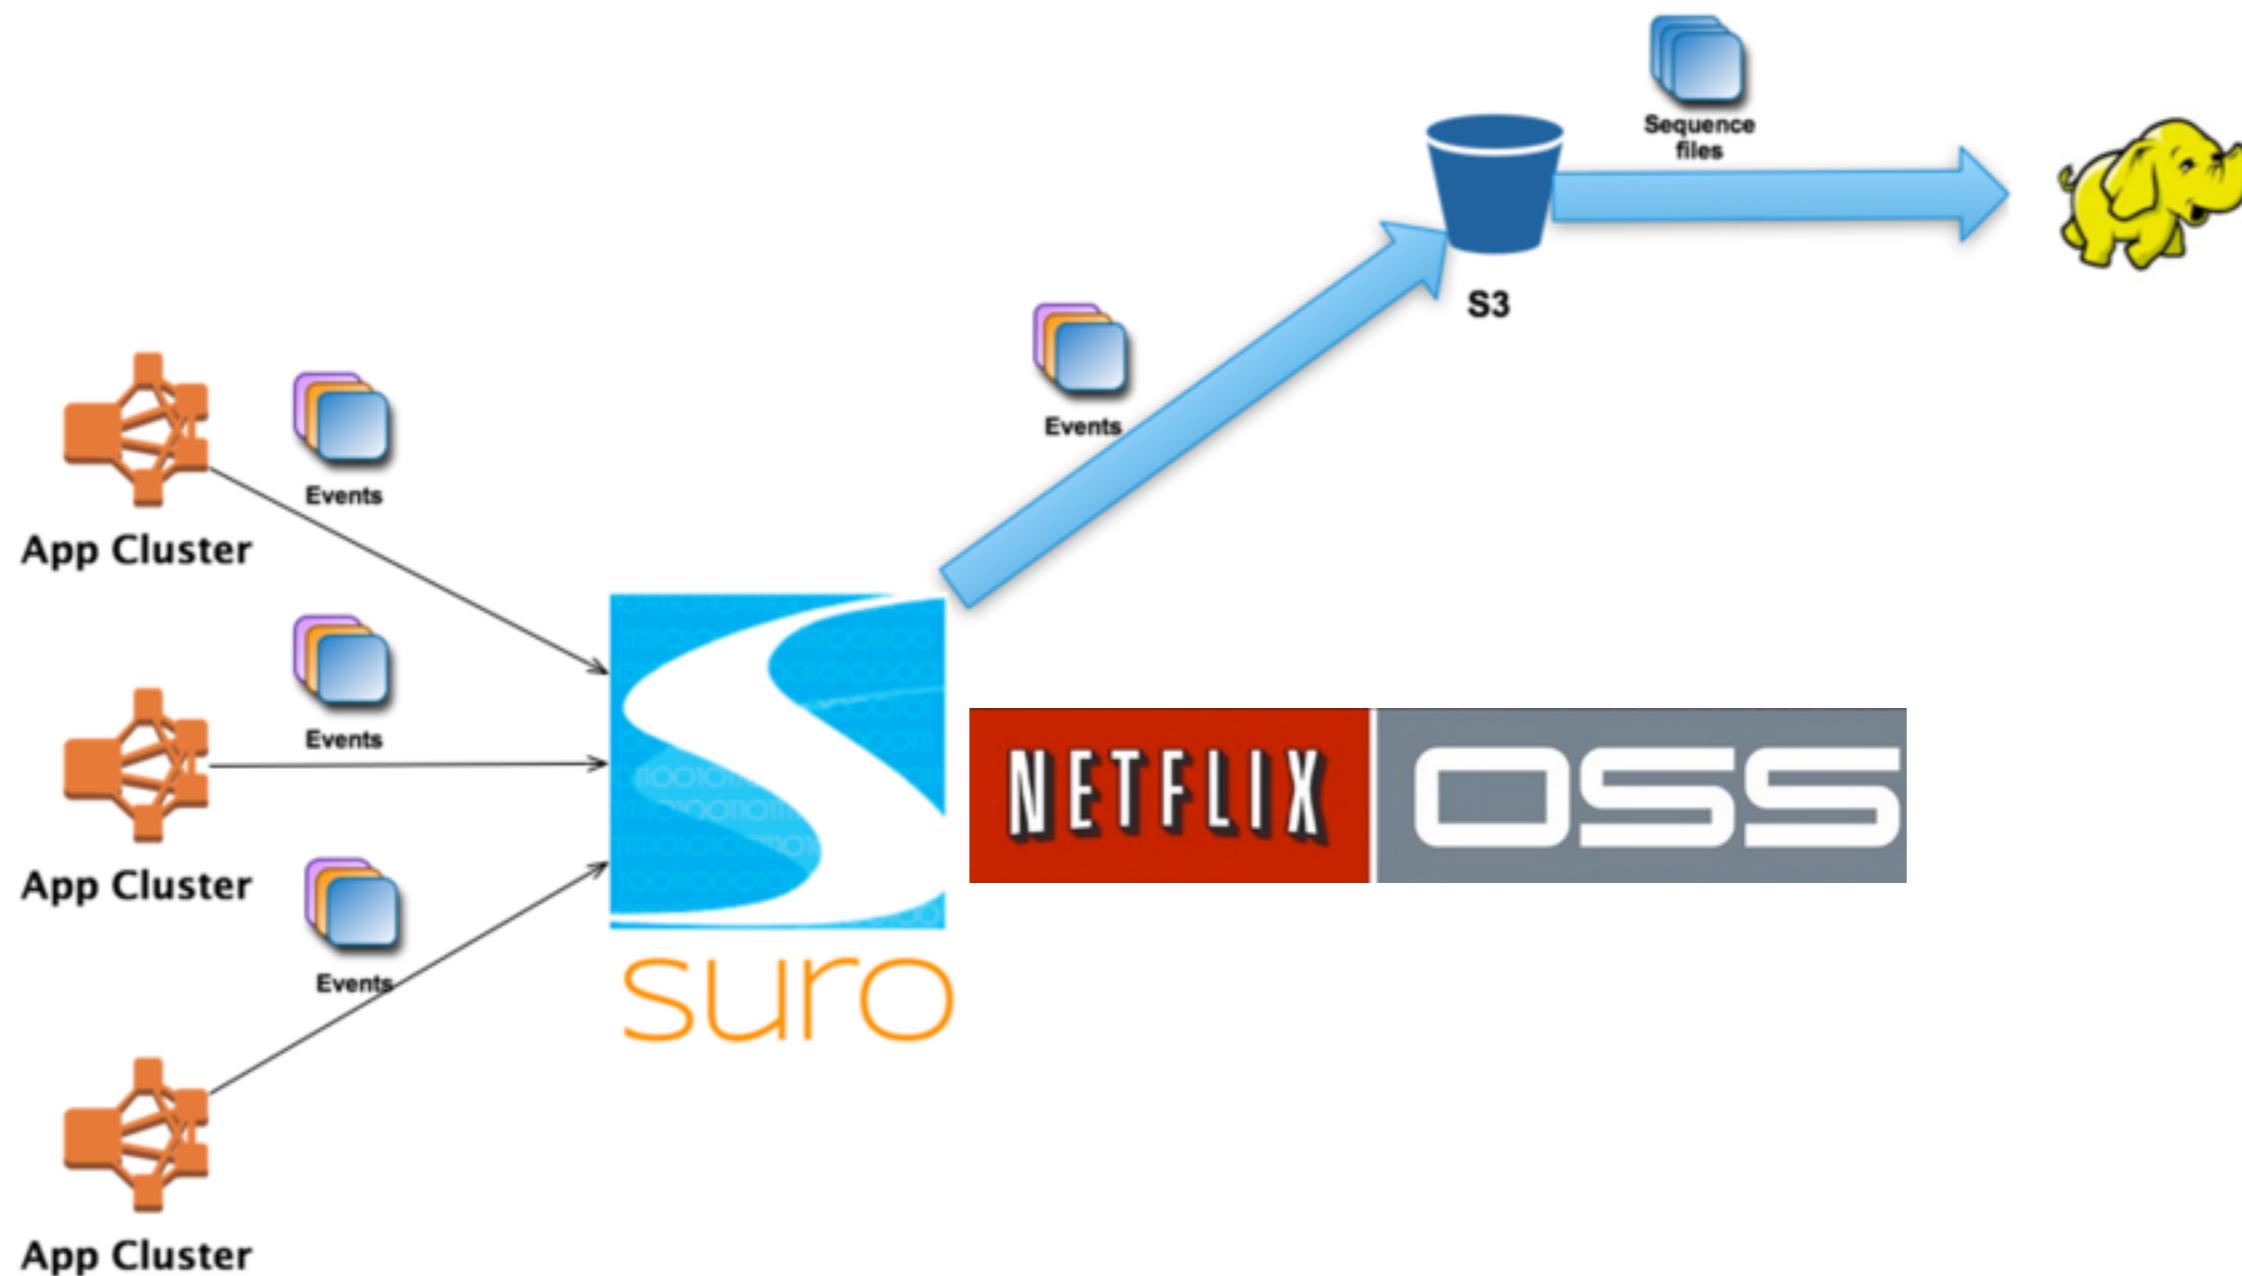
- Top N Query
- Trend
- Distribution

Non-Functional Requirements

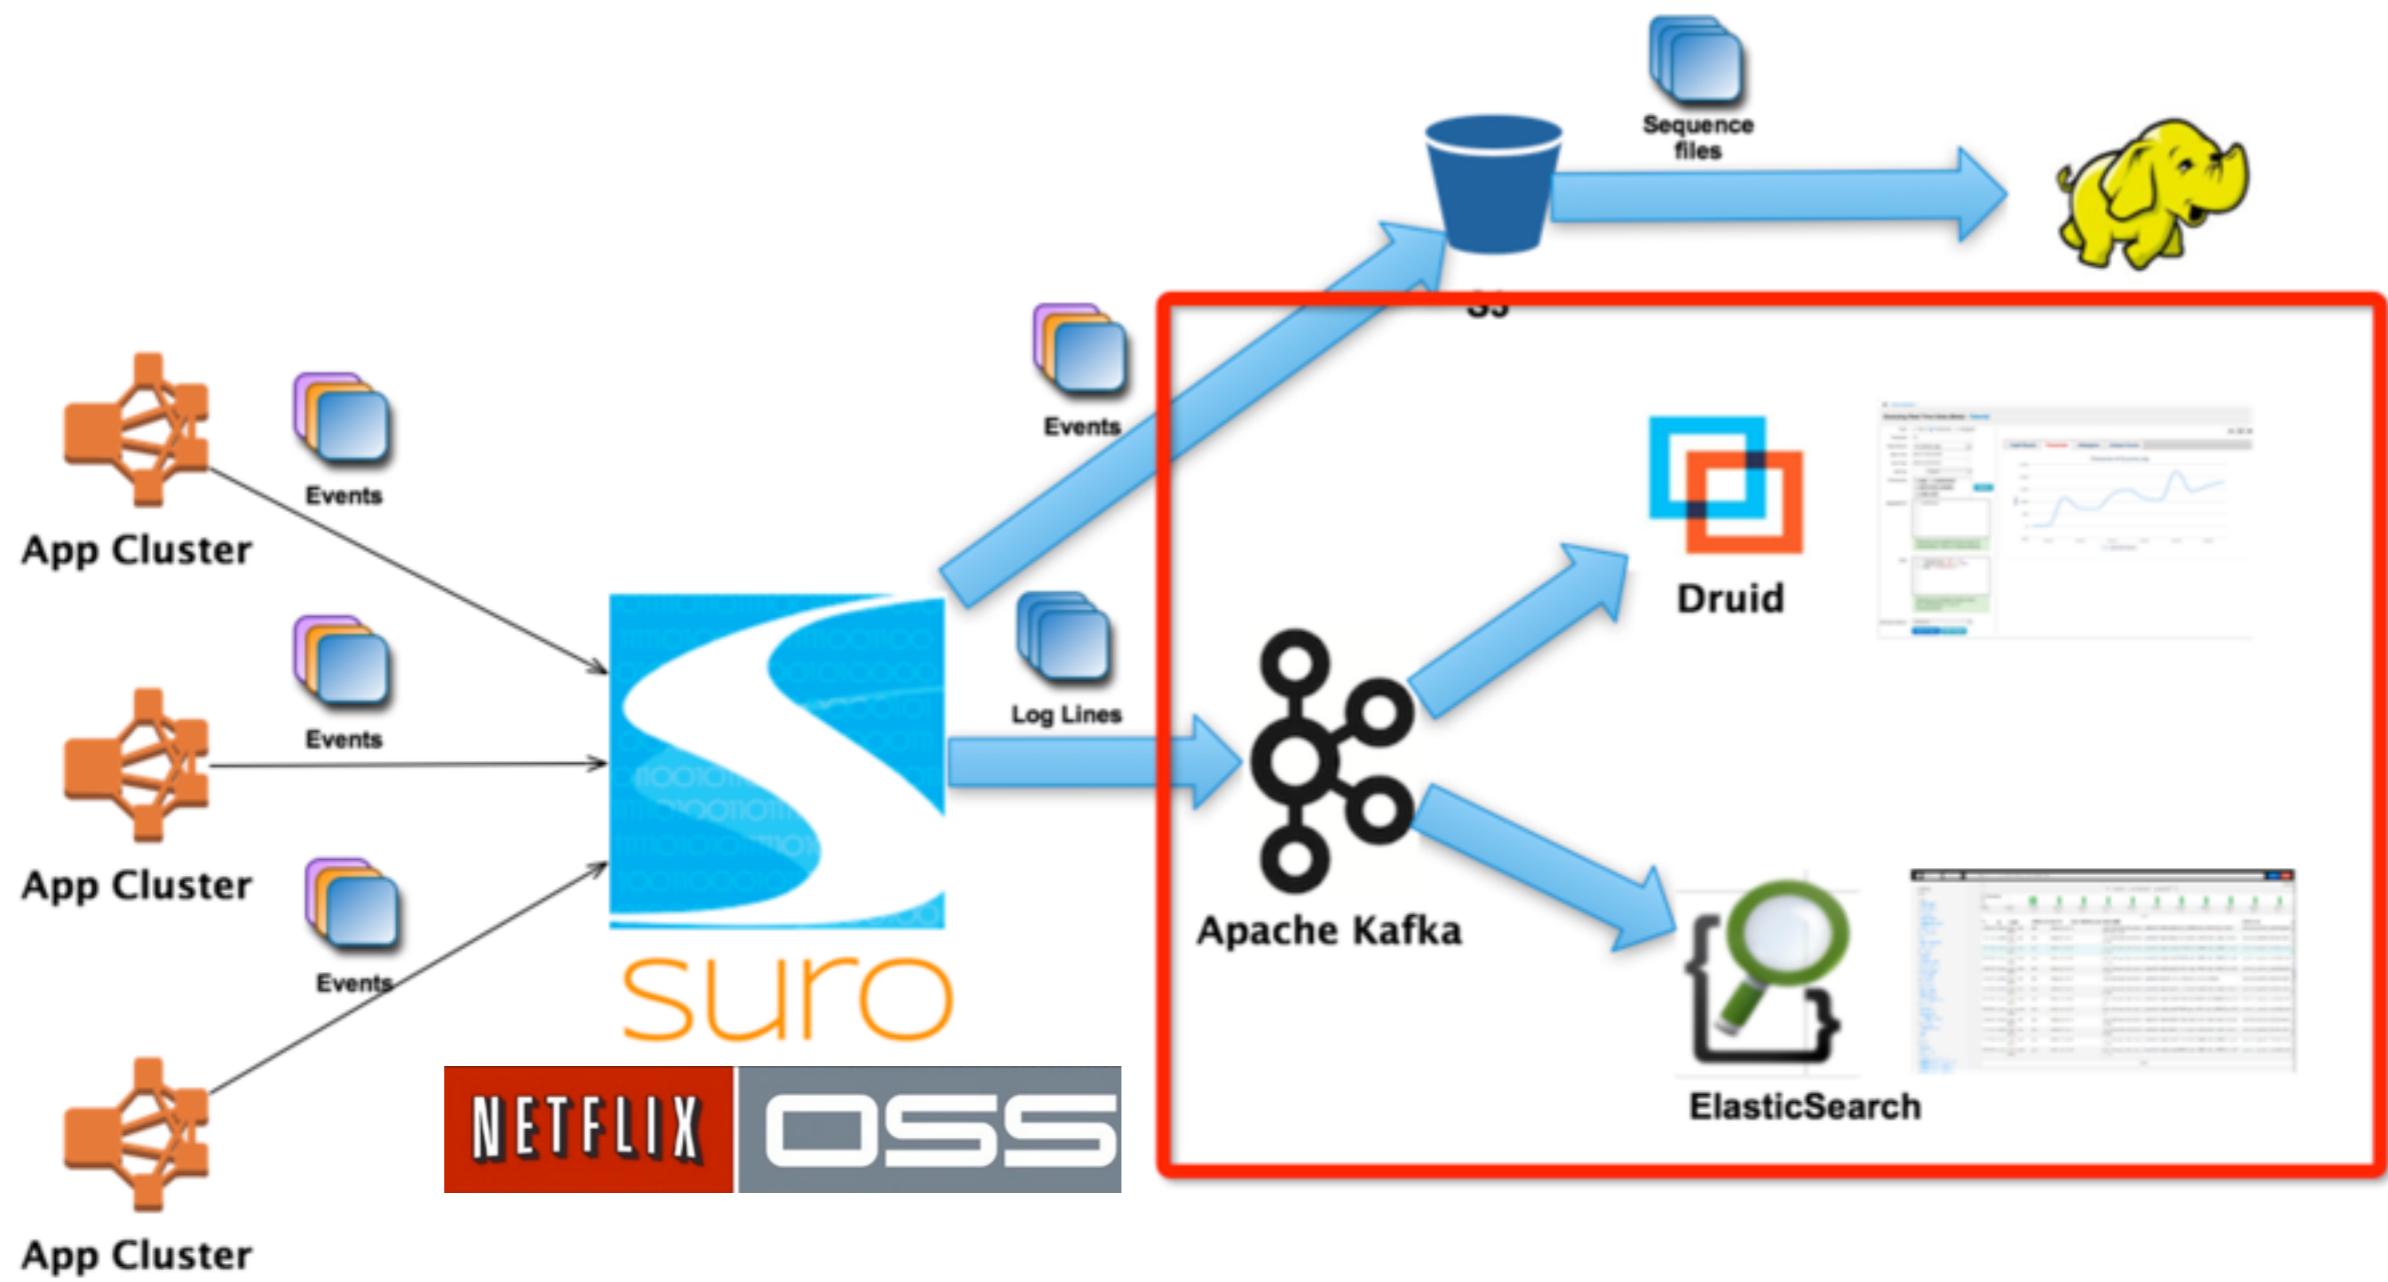
- Interactive (response within seconds)
- Quickly locates the right log events
- Minimal programming effort

**It's All about Extracting Small Data
Out of Big Data**





NETFLIX®





Now Back to the Use Case



Intelligent Alerts

Investigate the Errors

Start with finding potential outliers among cluster instances:
[Top 10 Errors grouped by Application, Logging Class, Host Name, and Line Number Log Summary Doc](#)

You should also check out dominant errors:

1. Find dominant errors by removing "hostname" from Dimensions box, and click on the button "Submit Query" below. You
2. Click on any text in the error table cell (they are links), and you'll see error details

Details of Error Surge

Contacts [Contacts](#)
About
Automated Log [Log Summary Documentation](#) [About Automated Alerts](#)
Summary Alerts
NAC [RECS PRECOMPUTE](#)
Trend Explorer [Error Trend that Caused This Alert \(RTExplorer's Documentation\)](#)

Trigger Graph

Time	errorCount per minute
08:41	1
08:42	3
08:43	4
08:44	7
08:45	9
08:46	10
08:47	9.5
08:48	7
08:49	4
08:50	2
08:51	1

Time



Guided Debugging in the Right Context

rtexplorer 1.1 us-east-1 prod

rtexplorer /

Querying Real Time Data (Beta) - Tutorial

Type: Top N Timeseries Histogram

Threshold: 10

Aggregations: `[[newCount]]`

Parsing failed: Cannot read property 'dataSource' of undefined

Filter: `((severity='40') or (severity='20')) and (app='NetflixWeb')`

Parsing successful (input size: 57 characters, 5.2e+3 characters/s)

Selected Metrics: `[[newCount]]`

Submit Query Show Query

Please wait. This may take a minute or two.

Copyright Netflix.com

at.gu C4_Q_DFA.ppt Mu01.ppt igw-0014401.ppt thomas97b (1).ppt Show All



Guided Debugging in the Right Context

rtexplorer 1.1 us-east-1 prod

rtexplorer /

Querying Real Time Data (Beta) - Tutorial

Type: Top N Timeseries Histogram

Threshold: 10

Aggregations: `[[newCount]]`

Parsing failed: Cannot read property 'dataSource' of undefined

Filter: `((severity='40') or (severity='20')) and (app='NetflixWeb')`

Parsing successful (input size: 57 characters, 5.2e+3 characters/s)

Selected Metrics: `[[newCount]]`

Submit Query Show Query

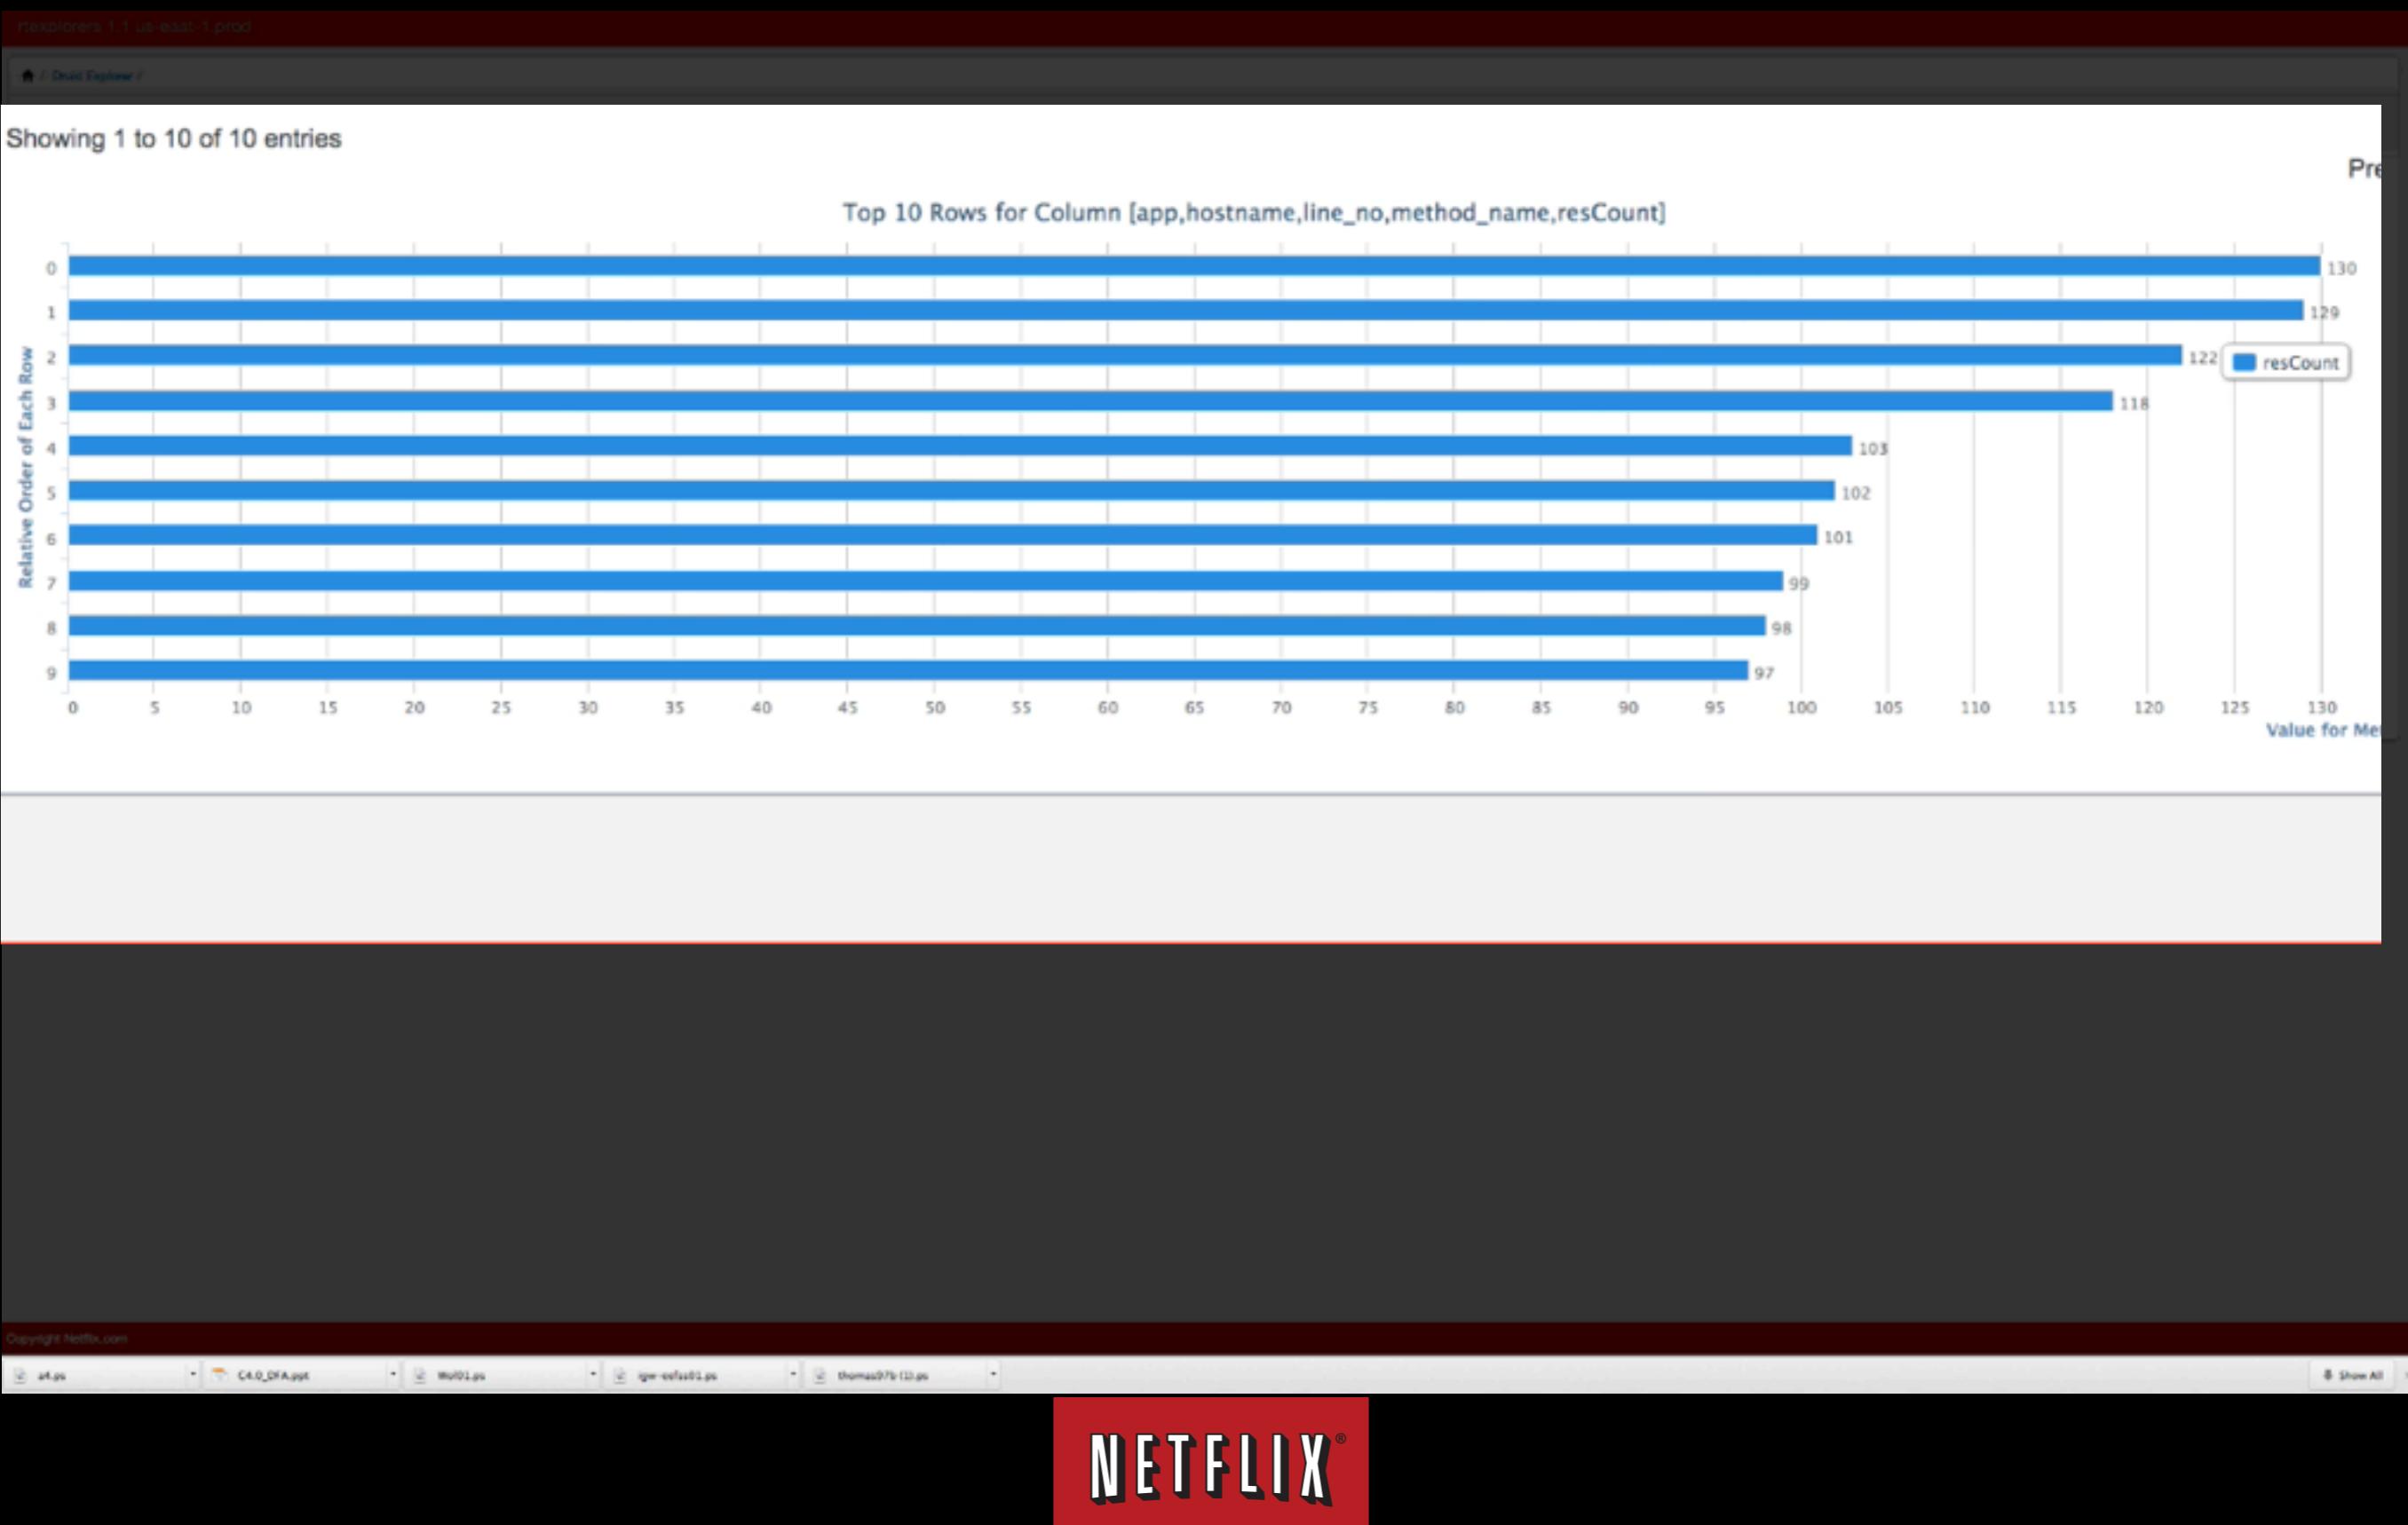
Please wait. This may take a minute or two.

Copyright Netflix.com

at.gu C4_Q_DFA.ppt Mu01.ppt igw-0014401.ppt thomas97b (1).ppt Show All



Guided Debugging in the Right Context



Guided Debugging in the Right Context

rtexplorer 1.1 us-east-1.prod

/ Druid Explorer /

Querying Real Time Data (Beta) - Tutorial

Type: Top N Timeseries Histogram

Threshold: 10

Data Source: nf_errors_log

Start Time: 2013-09-04 12:00

End Time: 2013-09-04 13:11

Shift By: Days

Dimensions: app, hostname, METHOD_NAME, LINE_NO

Aggregations: resCount

Parsing successful (input size: 8 characters, 2.7e+3 characters/s)

Filter: ((severity='40') or (severity='30')) and (app='MerchWeb')

Parsing successful (input size: 57 characters, 2.8e+4 characters/s)

Selected Metrics: resCount

Submit Query Show Query

TopN Result Timeseries Histogram Unique Count

Show 10 entries

app	hostname	line_no	method_name	resCount	Time
MerchWeb	merchweb-usca-live-i-801523ee	317	executeWithResponse	130	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-961523fb	317	executeWithResponse	129	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-801523ee	164	getProfiles	122	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-961523fb	164	getProfiles	118	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-1ae7137d	291	getNavigationItemsFromMap	103	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-9c1523f2	317	executeWithResponse	102	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-8a1523e4	291	getNavigationItemsFromMap	101	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-801523ee	291	getNavigationItemsFromMap	99	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-8a1523e4	317	executeWithResponse	98	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-9c1523f2	164	getProfiles	97	2013-09-04T19:59:00/2013-09-04T20:12:00

Showing 1 to 10 of 10 entries

Previous Next

Top 10 Rows for Column [app,hostname,line_no,method_name,resCount]

Relative Order of Each Row

Value for Metric resCount

Guided Debugging in the Right Context

rtexplorer 1.1 us-east-1.prod

/ Druid Explorer /

Querying Real Time Data (Beta) - Tutorial

Type: Top N Timeseries Histogram

Threshold: 10

Data Source: nf_errors_log

Start Time: 2013-09-04 12:00

End Time: 2013-09-04 13:11

Shift By: Days

Dimensions: app, hostname, METHOD_NAME, LINE_NO

Aggregations: resCount

Parsing successful (input size: 8 characters, 2.7e+3 characters/s)

Filter: ((severity='40') or (severity='30')) and (app='MerchWeb')

Parsing successful (input size: 57 characters, 2.8e+4 characters/s)

Selected Metrics: resCount

Submit Query Show Query

TopN Result Timeseries Histogram Unique Count

Show 10 entries

app	hostname	line_no	method_name	resCount	Time
MerchWeb	merchweb-usca-live-i-801523ee	317	executeWithResponse	130	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-961523fb	317	executeWithResponse	129	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-801523ee	164	getProfiles	122	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-961523fb	164	getProfiles	118	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-1ae7137d	291	getNavigationItemsFromMap	103	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-9c1523f2	317	executeWithResponse	102	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-8a1523e4	291	getNavigationItemsFromMap	101	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-801523ee	291	getNavigationItemsFromMap	99	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-8a1523e4	317	executeWithResponse	98	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	merchweb-usca-live-i-9c1523f2	164	getProfiles	97	2013-09-04T19:59:00/2013-09-04T20:12:00

Showing 1 to 10 of 10 entries

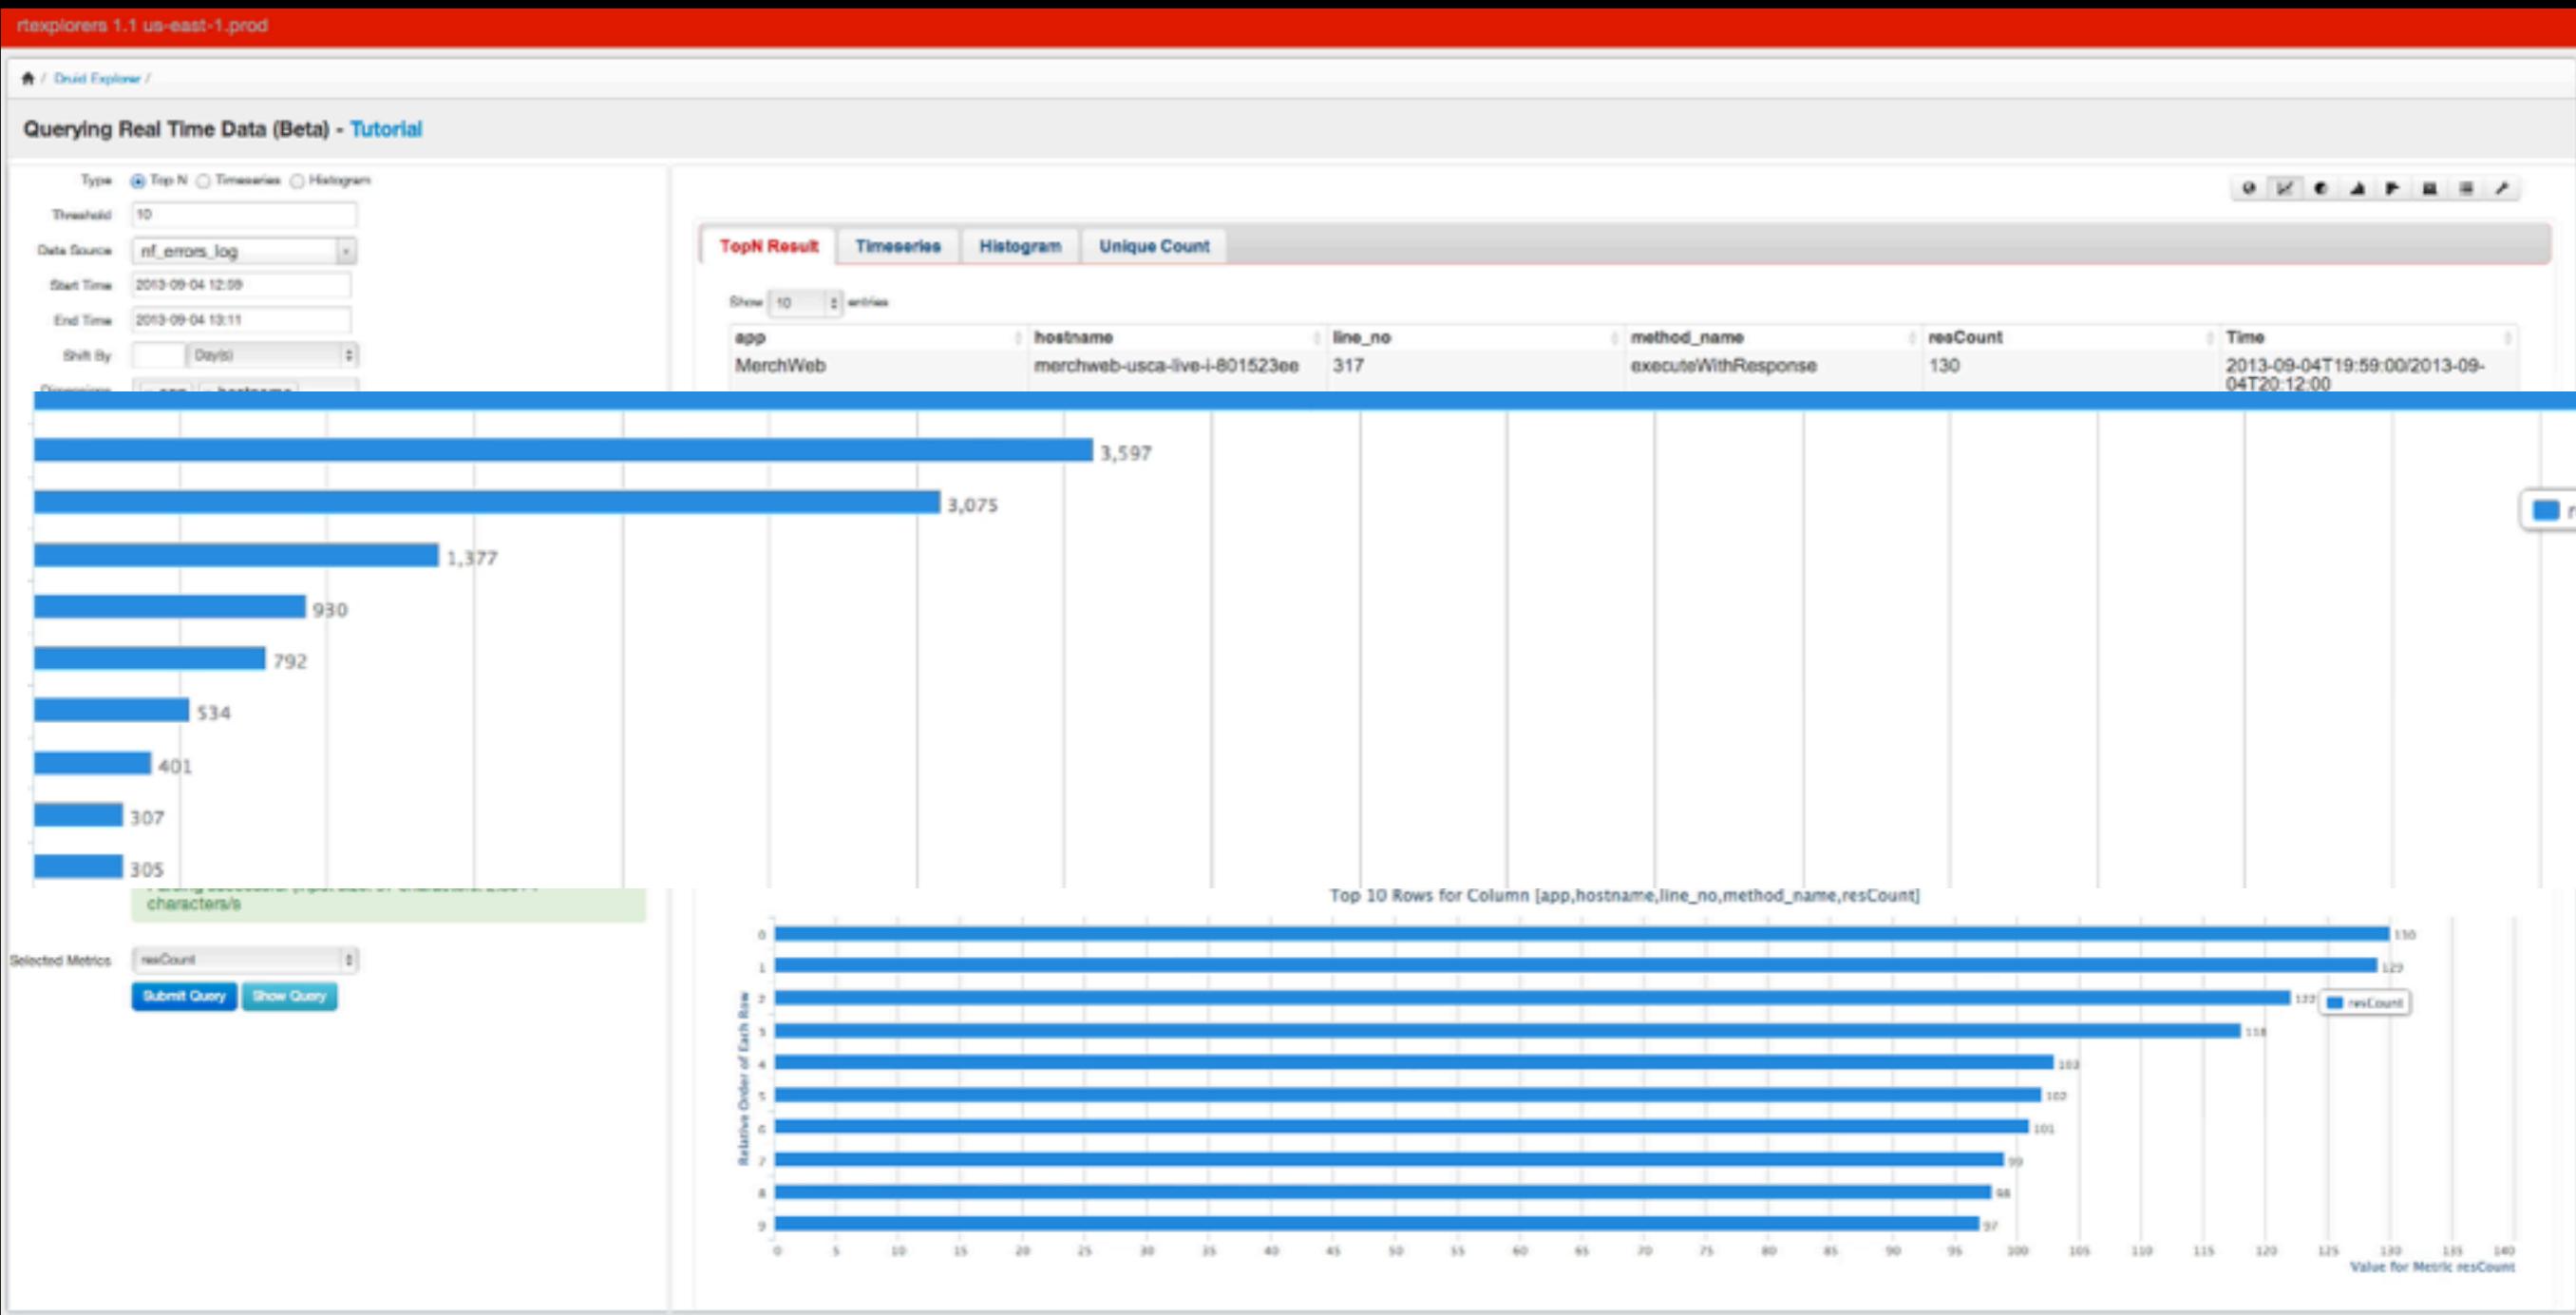
Previous Next

Top 10 Rows for Column [app,hostname,line_no,method_name,resCount]

Relative Order of Each Row

Value for Metric resCount

Guided Debugging in the Right Context



Guided Debugging in the Right Context

rtexplorers 1.1 us-east-1.prod

/ Druid Explorer /

Querying Real Time Data (Beta) - Tutorial

Type: Top N Timeseries Histogram

Threshold: 10

Data Source: nf_errors.log

Start Time: 2013-09-04 12:00

End Time: 2013-09-04 13:11

Shift By: [Day(s)]

Dimensions: app METHOD_NAME LINE_NO Metrics

Aggregations: resCount

Parsing successful (input size: 8 characters, 2.7e+3 characters/s)

Filter: `((severity='40') OR (severity='20')) AND (app='MerchWeb')`

Parsing successful (input size: 57 characters, 2.8e+4 characters/s)

Selected Metrics: resCount

Submit Query Show Query

TopN Result Timeseries Histogram Unique Count

Show: 10 entries

app	line_no	method_name	resCount	Time
MerchWeb	291	getNavigationItemsFromMap	8658	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	317	executeWithResponse	3597	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	164	getProfiles	3075	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	79	log	1377	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	488	setIdentityGracePeriod	930	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	284	_getVideoSimilar	792	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	339	run	534	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	631	run	401	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	71	run	307	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	63	getVideoHistoryItems	305	2013-09-04T19:59:00/2013-09-04T20:12:00

Showing 1 to 10 of 10 entries PreviousNext

Top 10 Rows for Column [app,line_no,method_name,resCount]

Relative Order of Each Row

Value for Metric resCount

Guided Debugging in the Right Context

rtexplorers 1.1 us-east-1.prod

/ Druid Explorer /

Querying Real Time Data (Beta) - Tutorial

Type: Top N Timeseries Histogram

Threshold: 10

Data Source: nf_errors.log

Start Time: 2013-09-04 12:00

End Time: 2013-09-04 13:11

Shift By: [Day(s)]

Dimensions: app METHOD_NAME LINE_NO Metrics

Aggregations: resCount

Parsing successful (input size: 8 characters, 2.7e+3 characters/s)

Filter: `((severity='40') OR (severity='20')) AND (app='MerchWeb')`

Parsing successful (input size: 57 characters, 2.8e+4 characters/s)

Selected Metrics: resCount

Submit Query Show Query

TopN Result				
Timeseries				
Histogram				
Unique Count				
Show: 10	entries			
app	line_no	method_name	resCount	Time
MerchWeb	291	getNavigationItemsFromMap	8658	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	317	executeWithResponse	3597	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	164	getProfiles	3075	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	79	log	1377	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	488	setIdentityGracePeriod	930	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	284	_getVideoSimilar	792	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	339	run	534	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	631	run	401	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	71	run	307	2013-09-04T19:59:00/2013-09-04T20:12:00
MerchWeb	63	getVideoHistoryItems	305	2013-09-04T19:59:00/2013-09-04T20:12:00

Showing 1 to 10 of 10 entries

Top 10 Rows for Column [app,line_no,method_name,resCount]

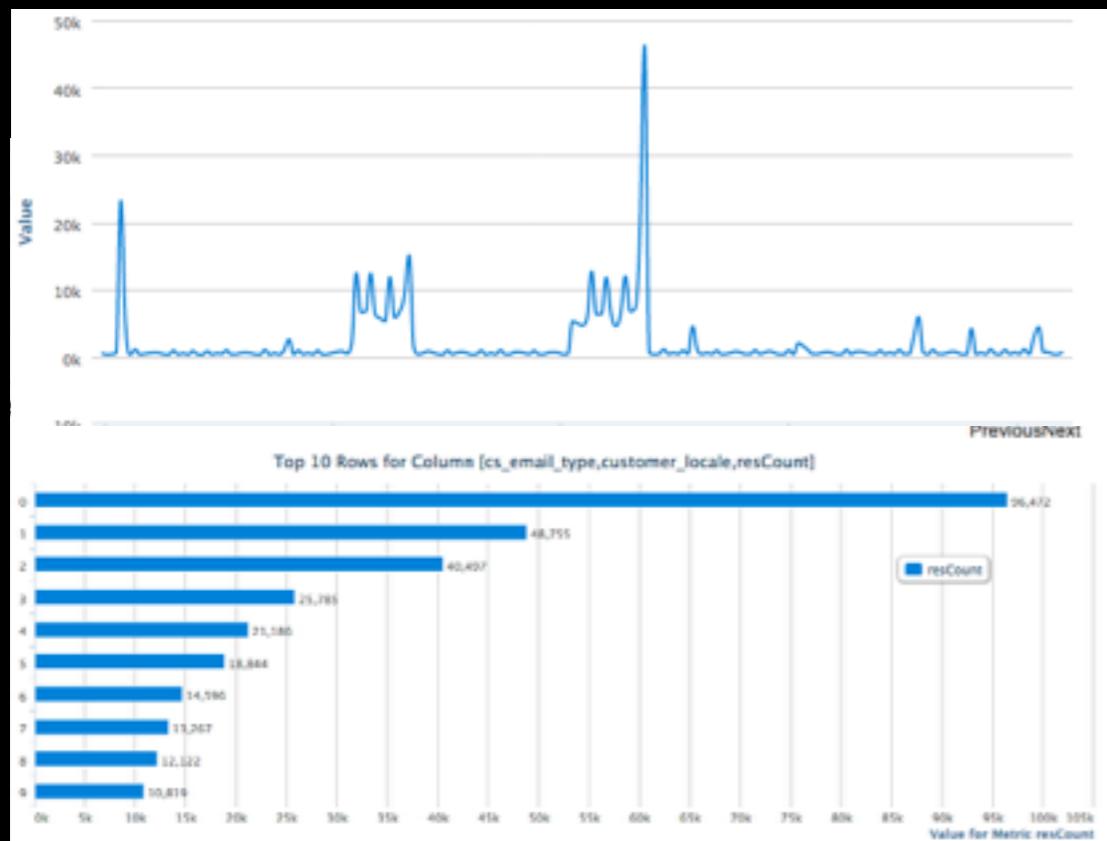
Relative Order of Each Row

Value for Metric resCount

A Useful Pattern



Aggregated Query -> Individual Query



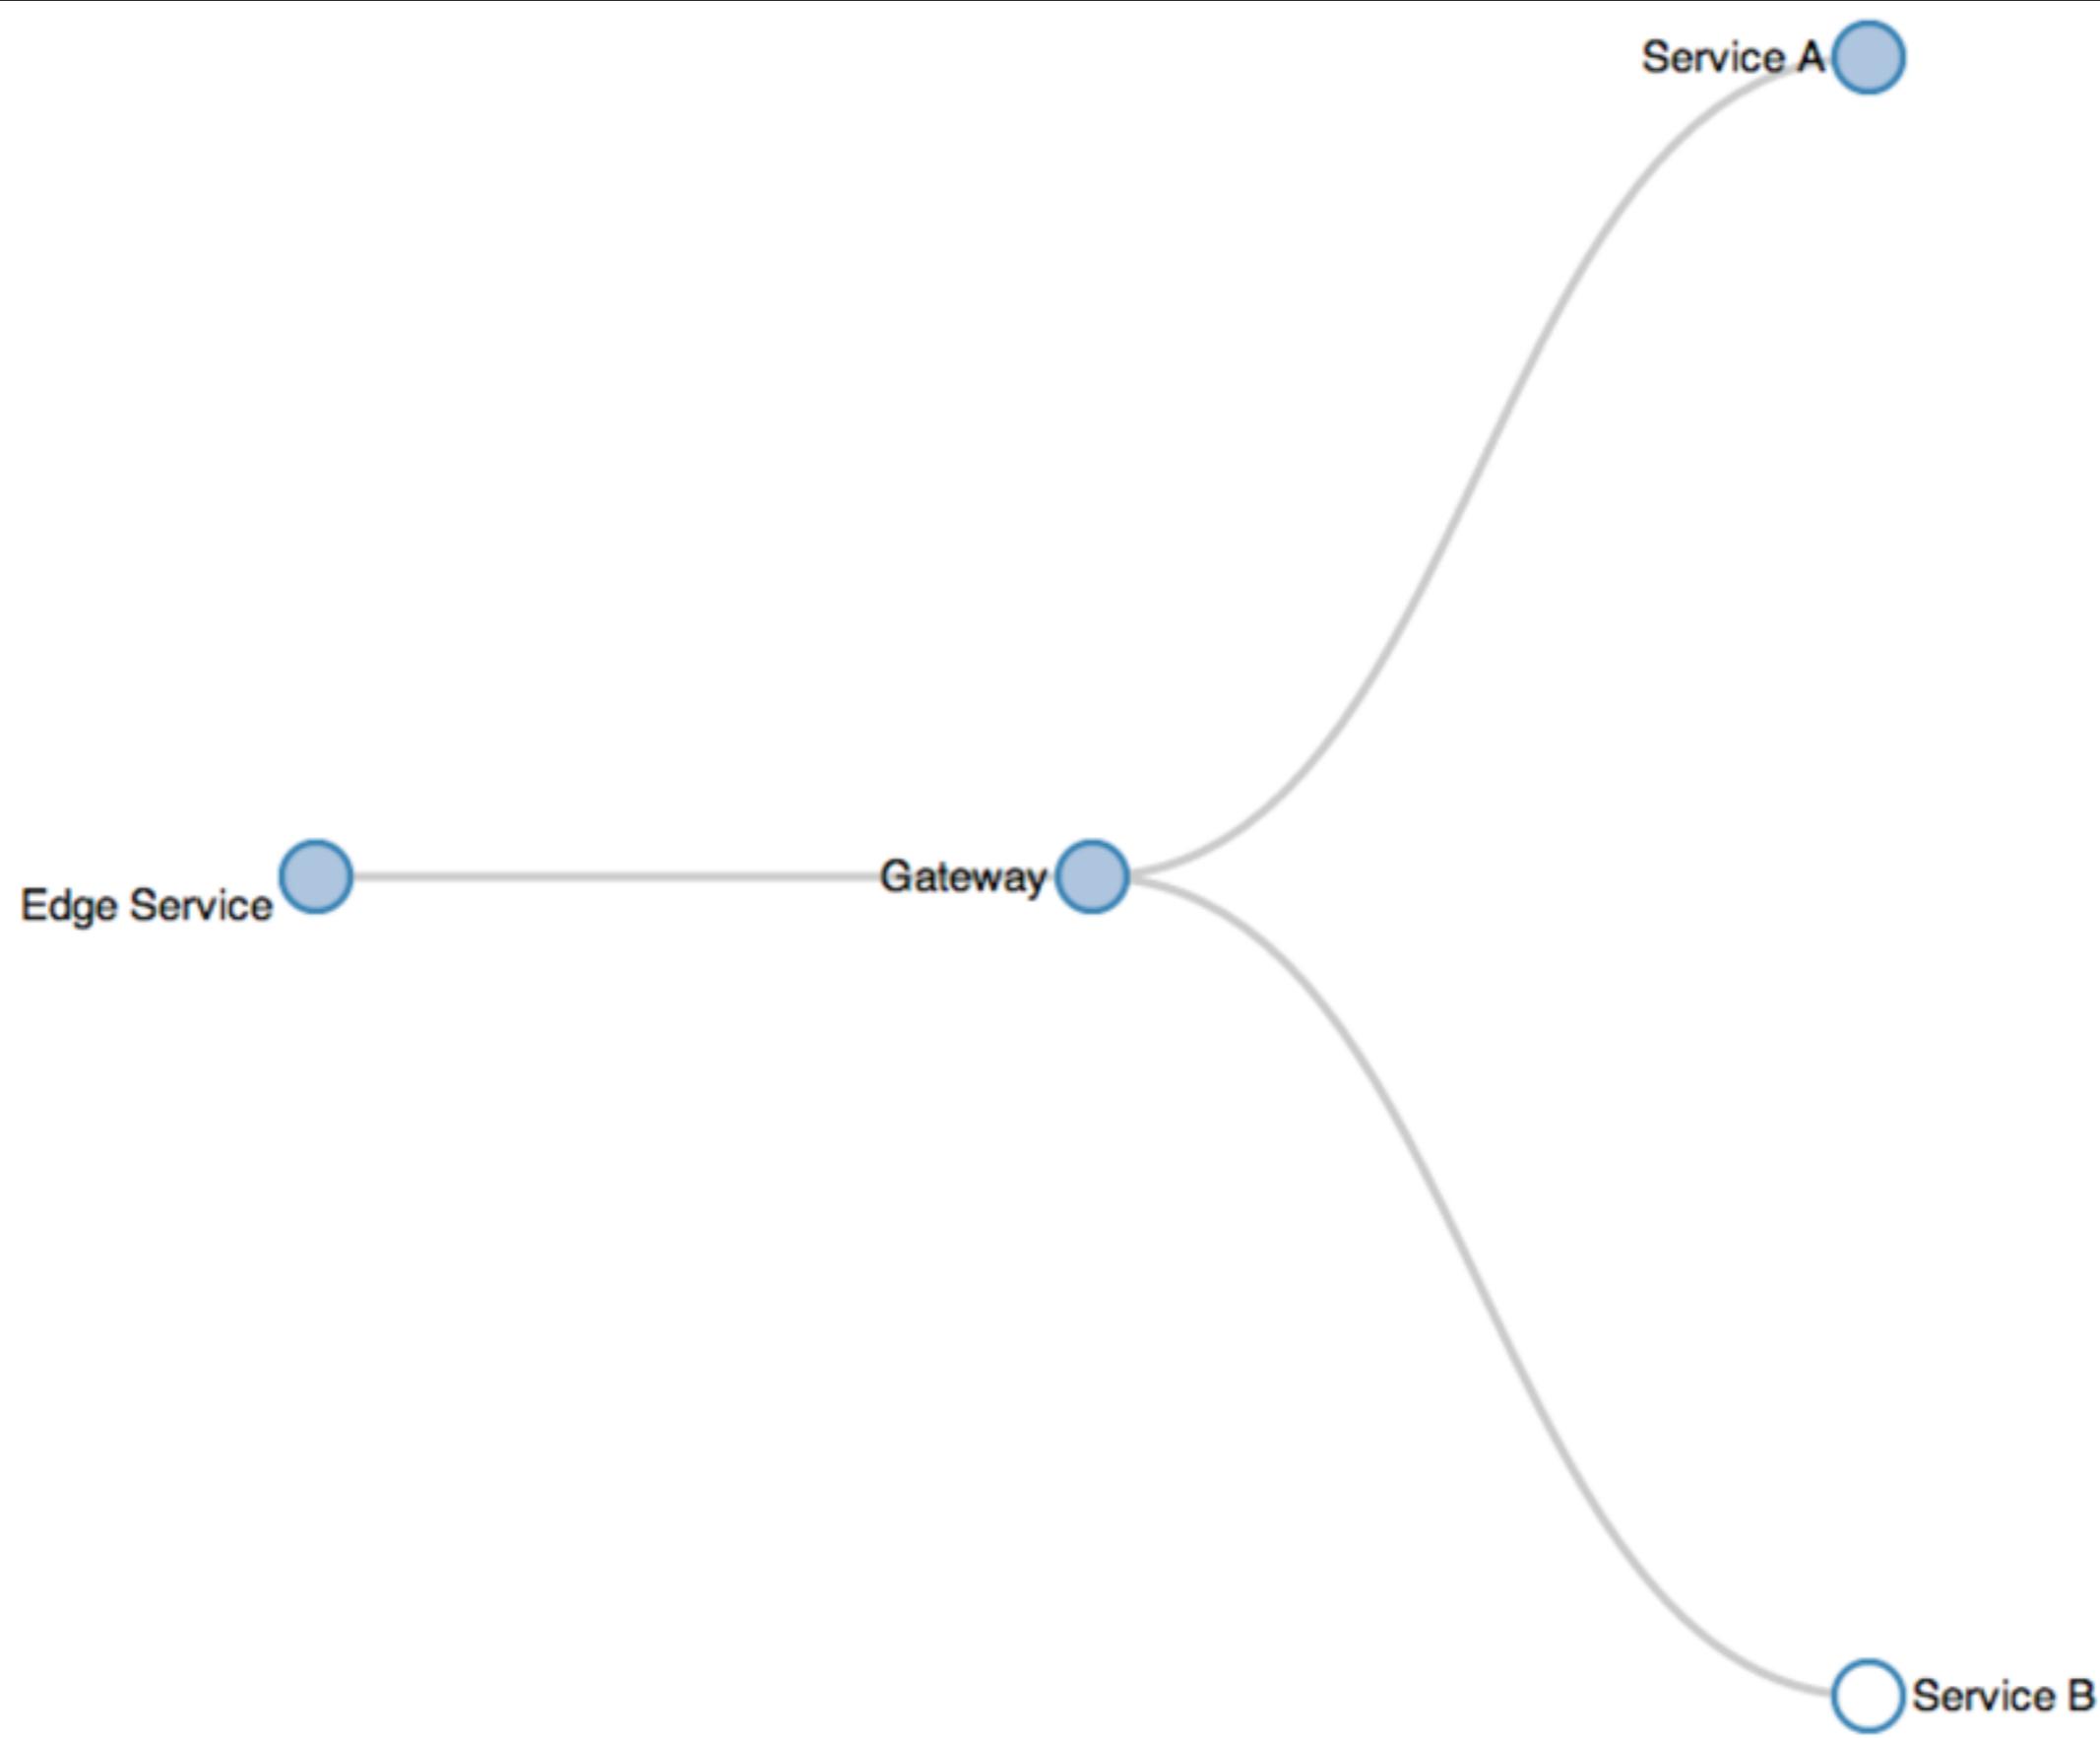
Examples

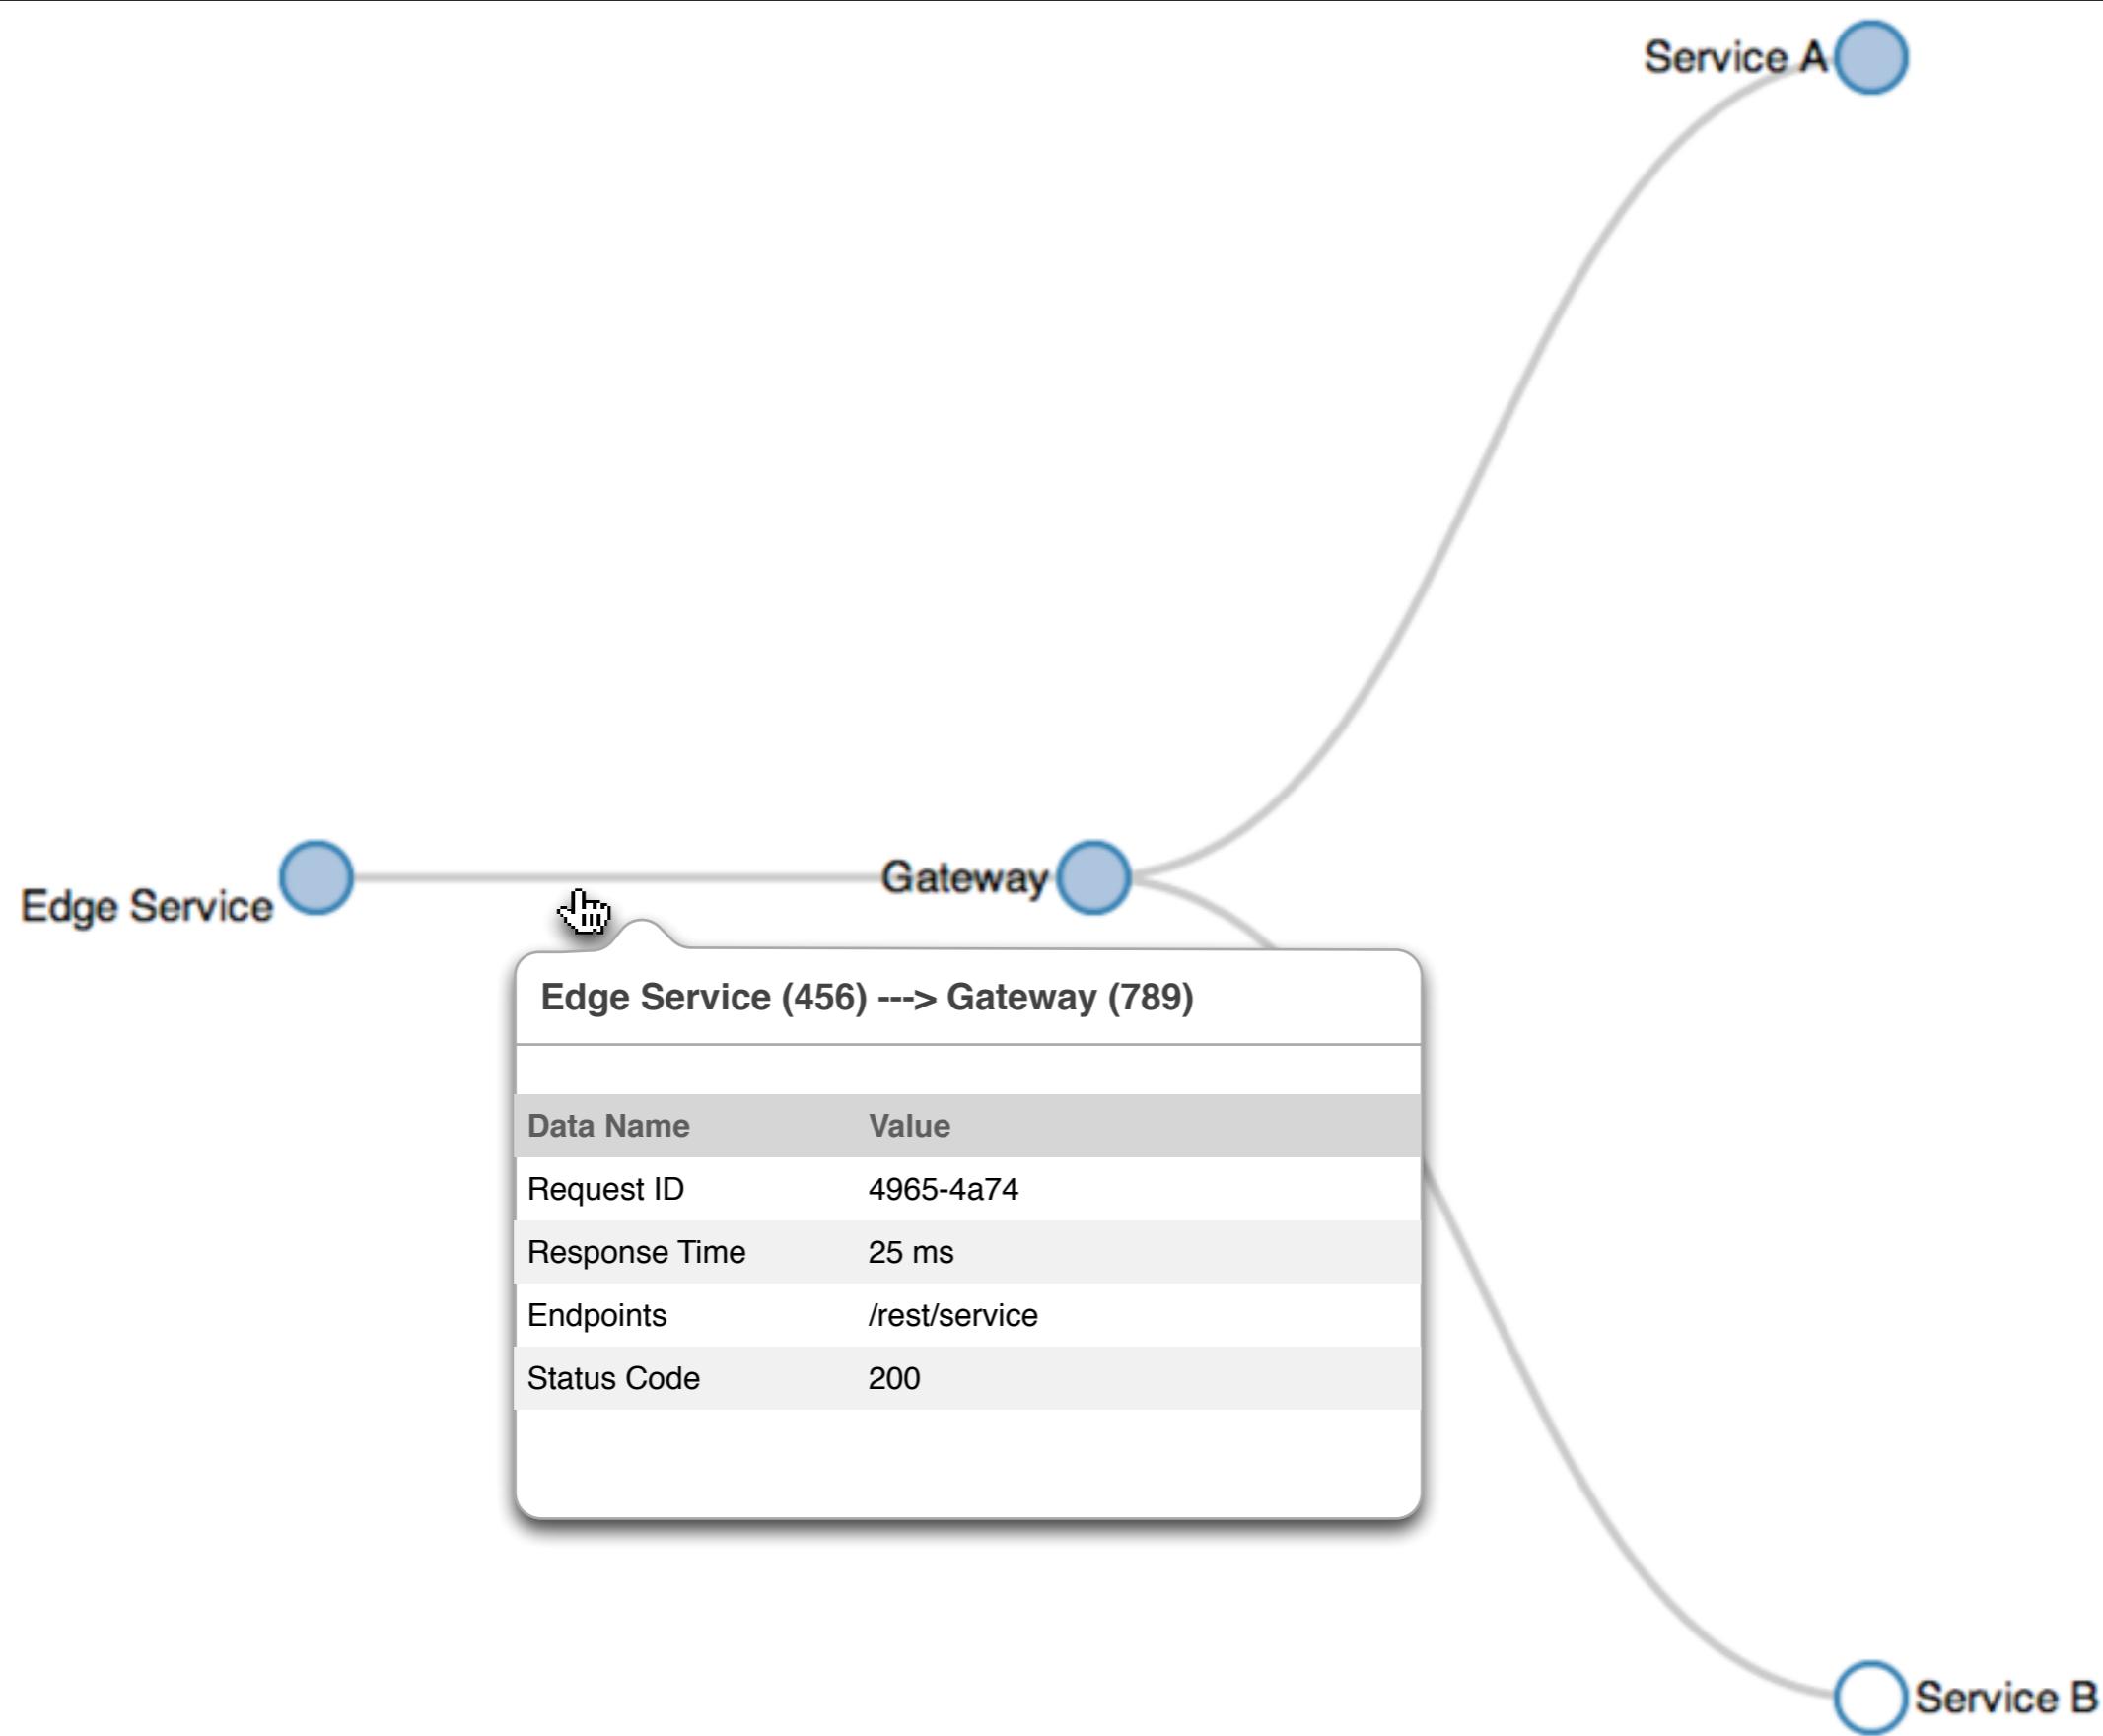
- S3 diagnostics
- Tracking email campaigns
- Request traces

Status:200

RequestId	Parent Id	Node Id	Service Name	Status
4965-4a74	0	123	Edge Service	200
4965-4a74	123	456	Gateway	200
4965-4a74	456	789	Service A	200
4965-4a74e	456	abc	Service B	200



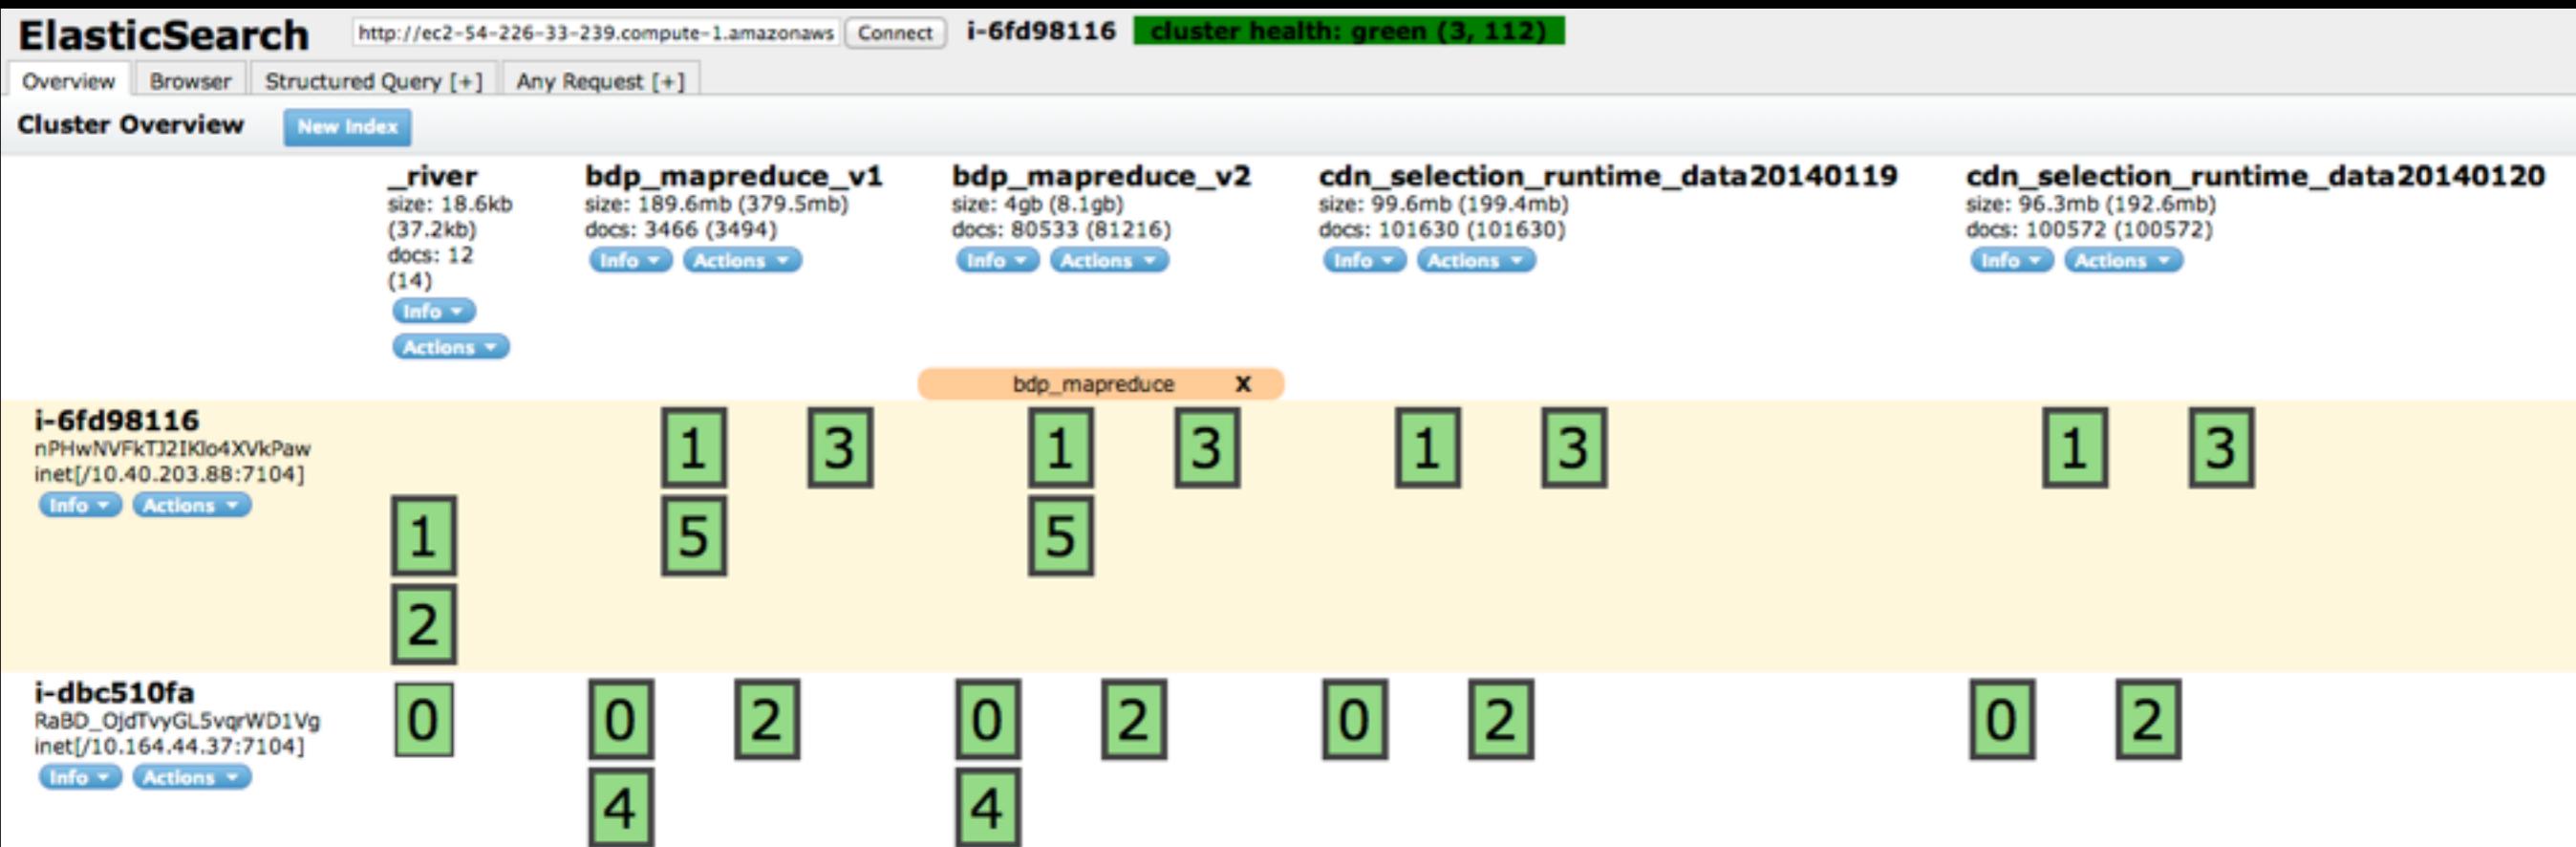




Why Elasticsearch?



Automatic Sharding and Replication





Flexible Schema



Flexible Schema

- Schemaless



Flexible Schema

- Schemaless
- Reasonable defaults



Nice Extension Model



Nice Extension Model

- Customizable REST Actions

Nice Extension Model

- Customizable REST Actions
- Site Plugins



Nice Extension Model

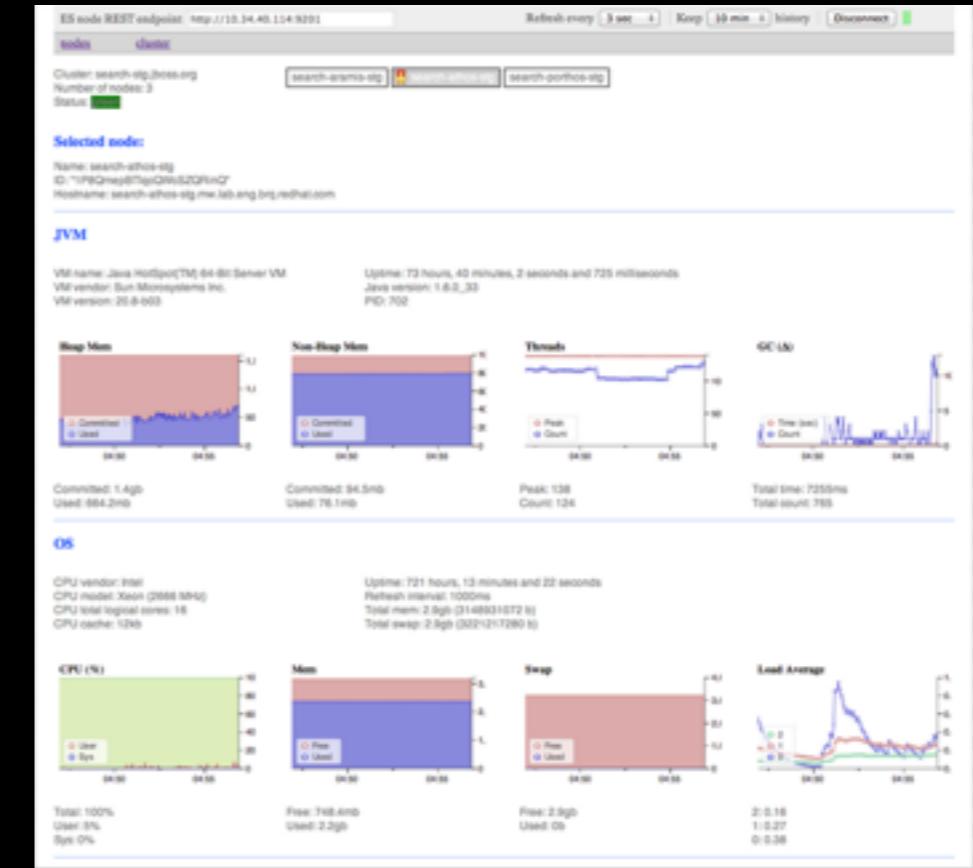
- Customizable REST Actions
- Site Plugins
- River Plugins



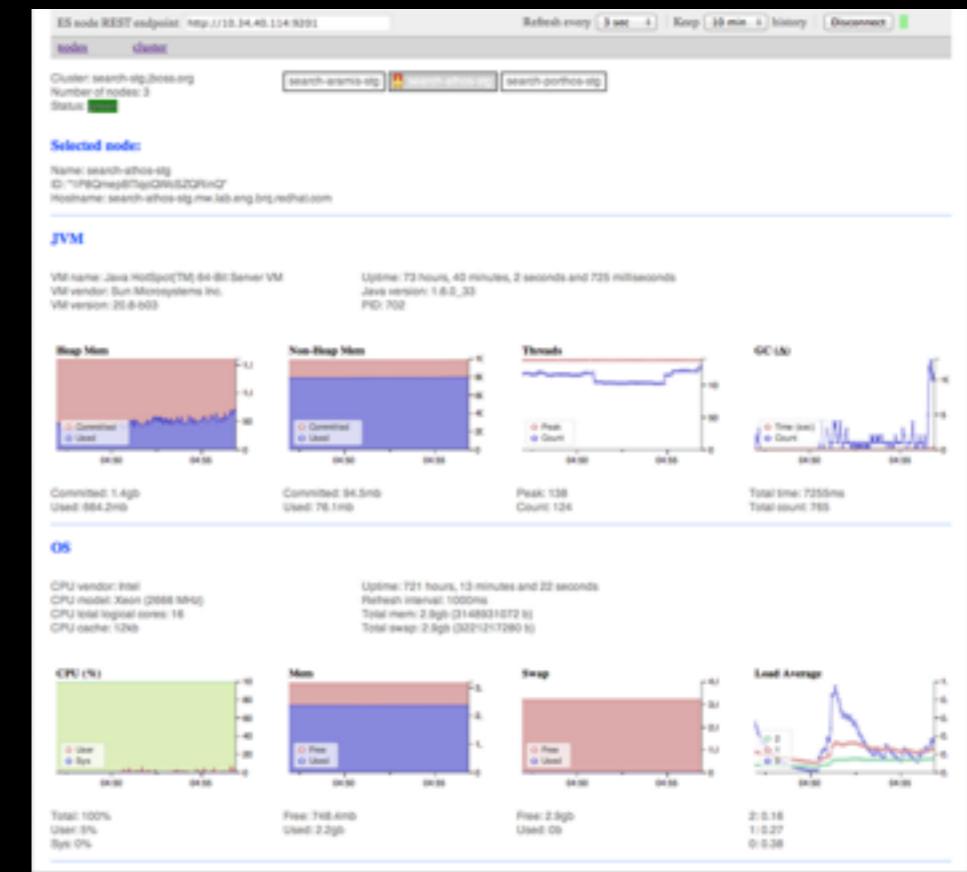
Nice Extension Model

- Customizable REST Actions
- Site Plugins
- River Plugins
- Discovery Module





Ecosystem - Plugins, Kibana



Tracking Service Deployments







NETFLIX®

Built by Netflix Monitoring Eng Team



NETFLIX®

Built by Netflix Monitoring Eng Team

Tracks History and Changes to Service Deployments



Built by Netflix Monitoring Eng Team

Tracks History and Changes to Service Deployments

Keeps Many Revisions



Built by Netflix Monitoring Eng Team

Tracks History and Changes to Service Deployments

Keeps Many Revisions

Tracks Dozens of Document Types



- Collection APIs
 - /aws
 - /aws/addresses
 - /aws/alarms
 - /aws/autoScalingGroups
 - /aws/buckets
 - /aws/databases
 - /aws/iamGroups
 - /aws/iamRoles
 - /aws/iamUsers
 - /aws/iamVirtualMFADevices
 - /aws/images
 - /aws/instances
 - /aws/launchConfigurations
 - /aws/loadBalancers
 - /aws/reservedInstances
 - /aws/scalingPolicies
 - /aws/securityGroups
 - /aws/snapshots
 - /aws/tags
 - /aws/volumes
 - /group
 - /group/autoScalingGroups
 - /view
 - /view/instances
 - /view/loadBalancerInstances
 - /view/simpleQueues



- Collection APIs
 - /aws
 - /aws/addresses
 - /aws/alarms
 - /aws/autoScalingGroups
 - /aws/buckets
 - /aws/databases
 - /aws/iamGroups
 - /aws/iamRoles
 - /aws/iamUsers
 - /aws/iamVirtualMFADevices
 - /aws/images
 - /aws/instances
 - /aws/launchConfigurations
 - /aws/loadBalancers
 - /aws/reservedInstances
 - /aws/scalingPolicies
 - /aws/securityGroups
 - /aws/snapshots
 - /aws/tags
 - /aws/volumes
 - /group
 - /group/autoScalingGroups
 - /view
 - /view/instances
 - /view/loadBalancerInstances
 - /view/simpleQueues

- General
 - Select Matrix Arguments
 - Modifier Matrix Arguments
 - _all
 - _at=<ms timestamp>
 - _callback=<name>
 - _diff[=<context>]
 - _expand
 - _limit=<num>
 - _live
 - _meta
 - _pp
 - _since=<ms timestamp>
 - _until=<ms timestamp>
 - _updated
 - Field Selectors



Why Elasticsearch?





Schemas may change at any time



Schemas may change at any time

Go schemaless





Users may search for any combination of fields



Users may search for any combination of fields

This is what search engine is designed for





Users often needs only a few fields



Users often needs only a few fields

Projection via “fields” query





Need range queries on date and revisions



Need range queries on date and revisions

Natively supported by Elasticsearch



Need range queries on date and revisions

Natively supported by Elasticsearch

Route by document ID



Running ES in Netflix



Operational Challenges



Operational Challenges

Back pressure when indexing

Operational Challenges

Back pressure when indexing

Diverse configurations and data

Operational Challenges

Back pressure when indexing

Diverse configurations and data

Dynamic flow of log events

Operational Challenges

Back pressure when indexing

Diverse configurations and data

Dynamic flow of log events

Needs extensive monitoring and alerting

Operational Challenges

Back pressure when indexing

Diverse configurations and data

Dynamic flow of log events

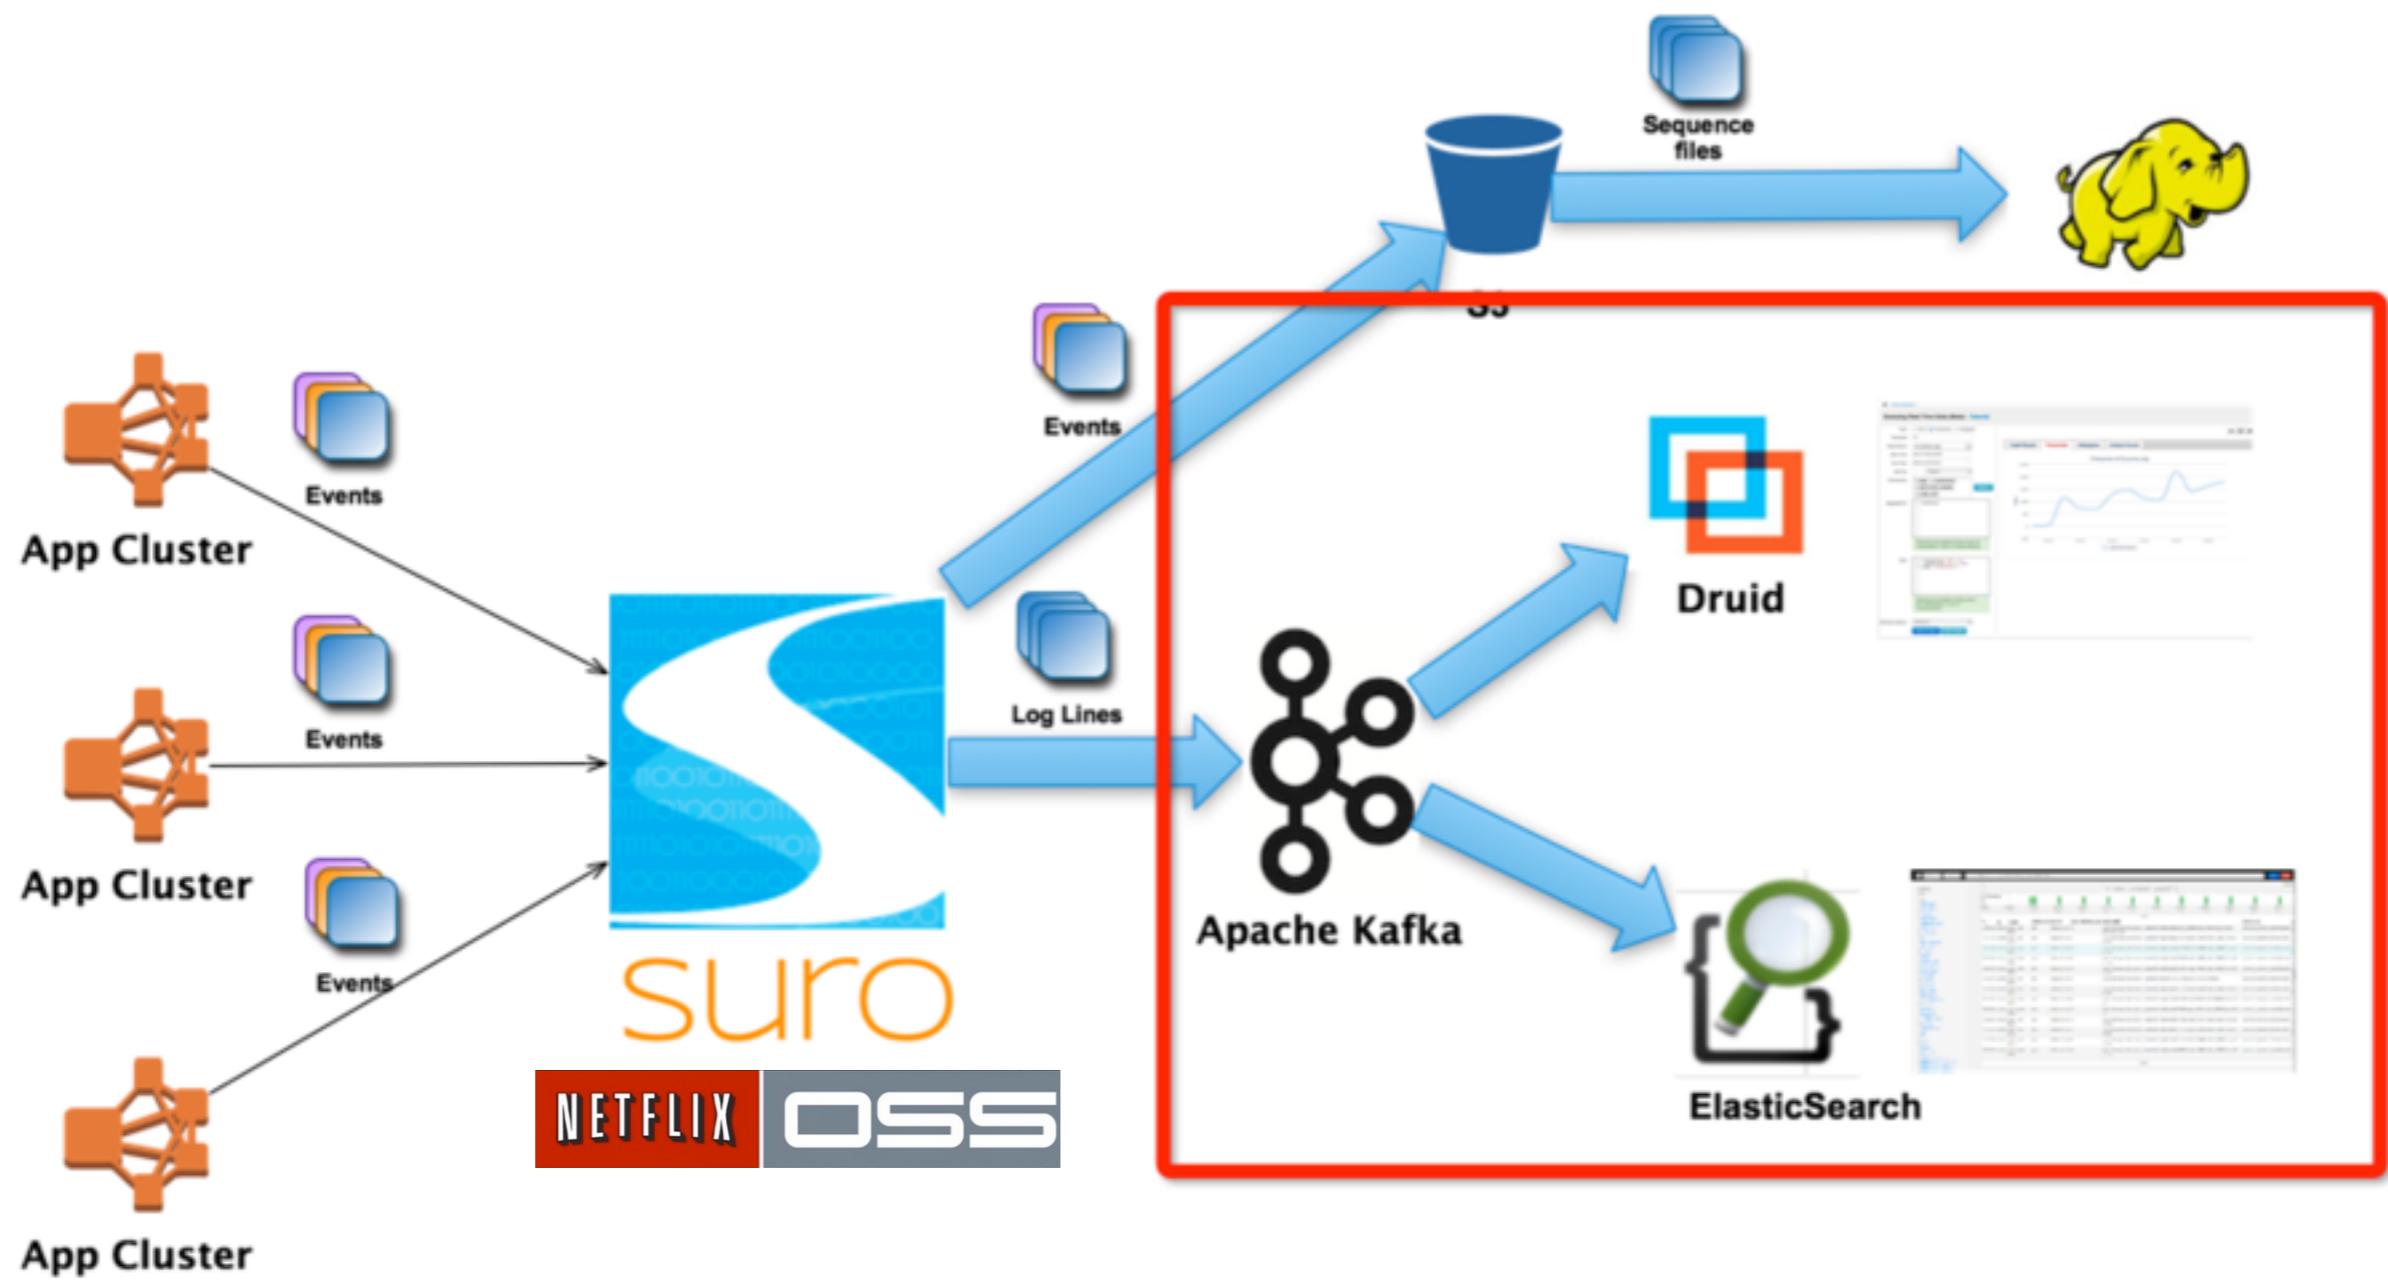
Needs extensive monitoring and alerting

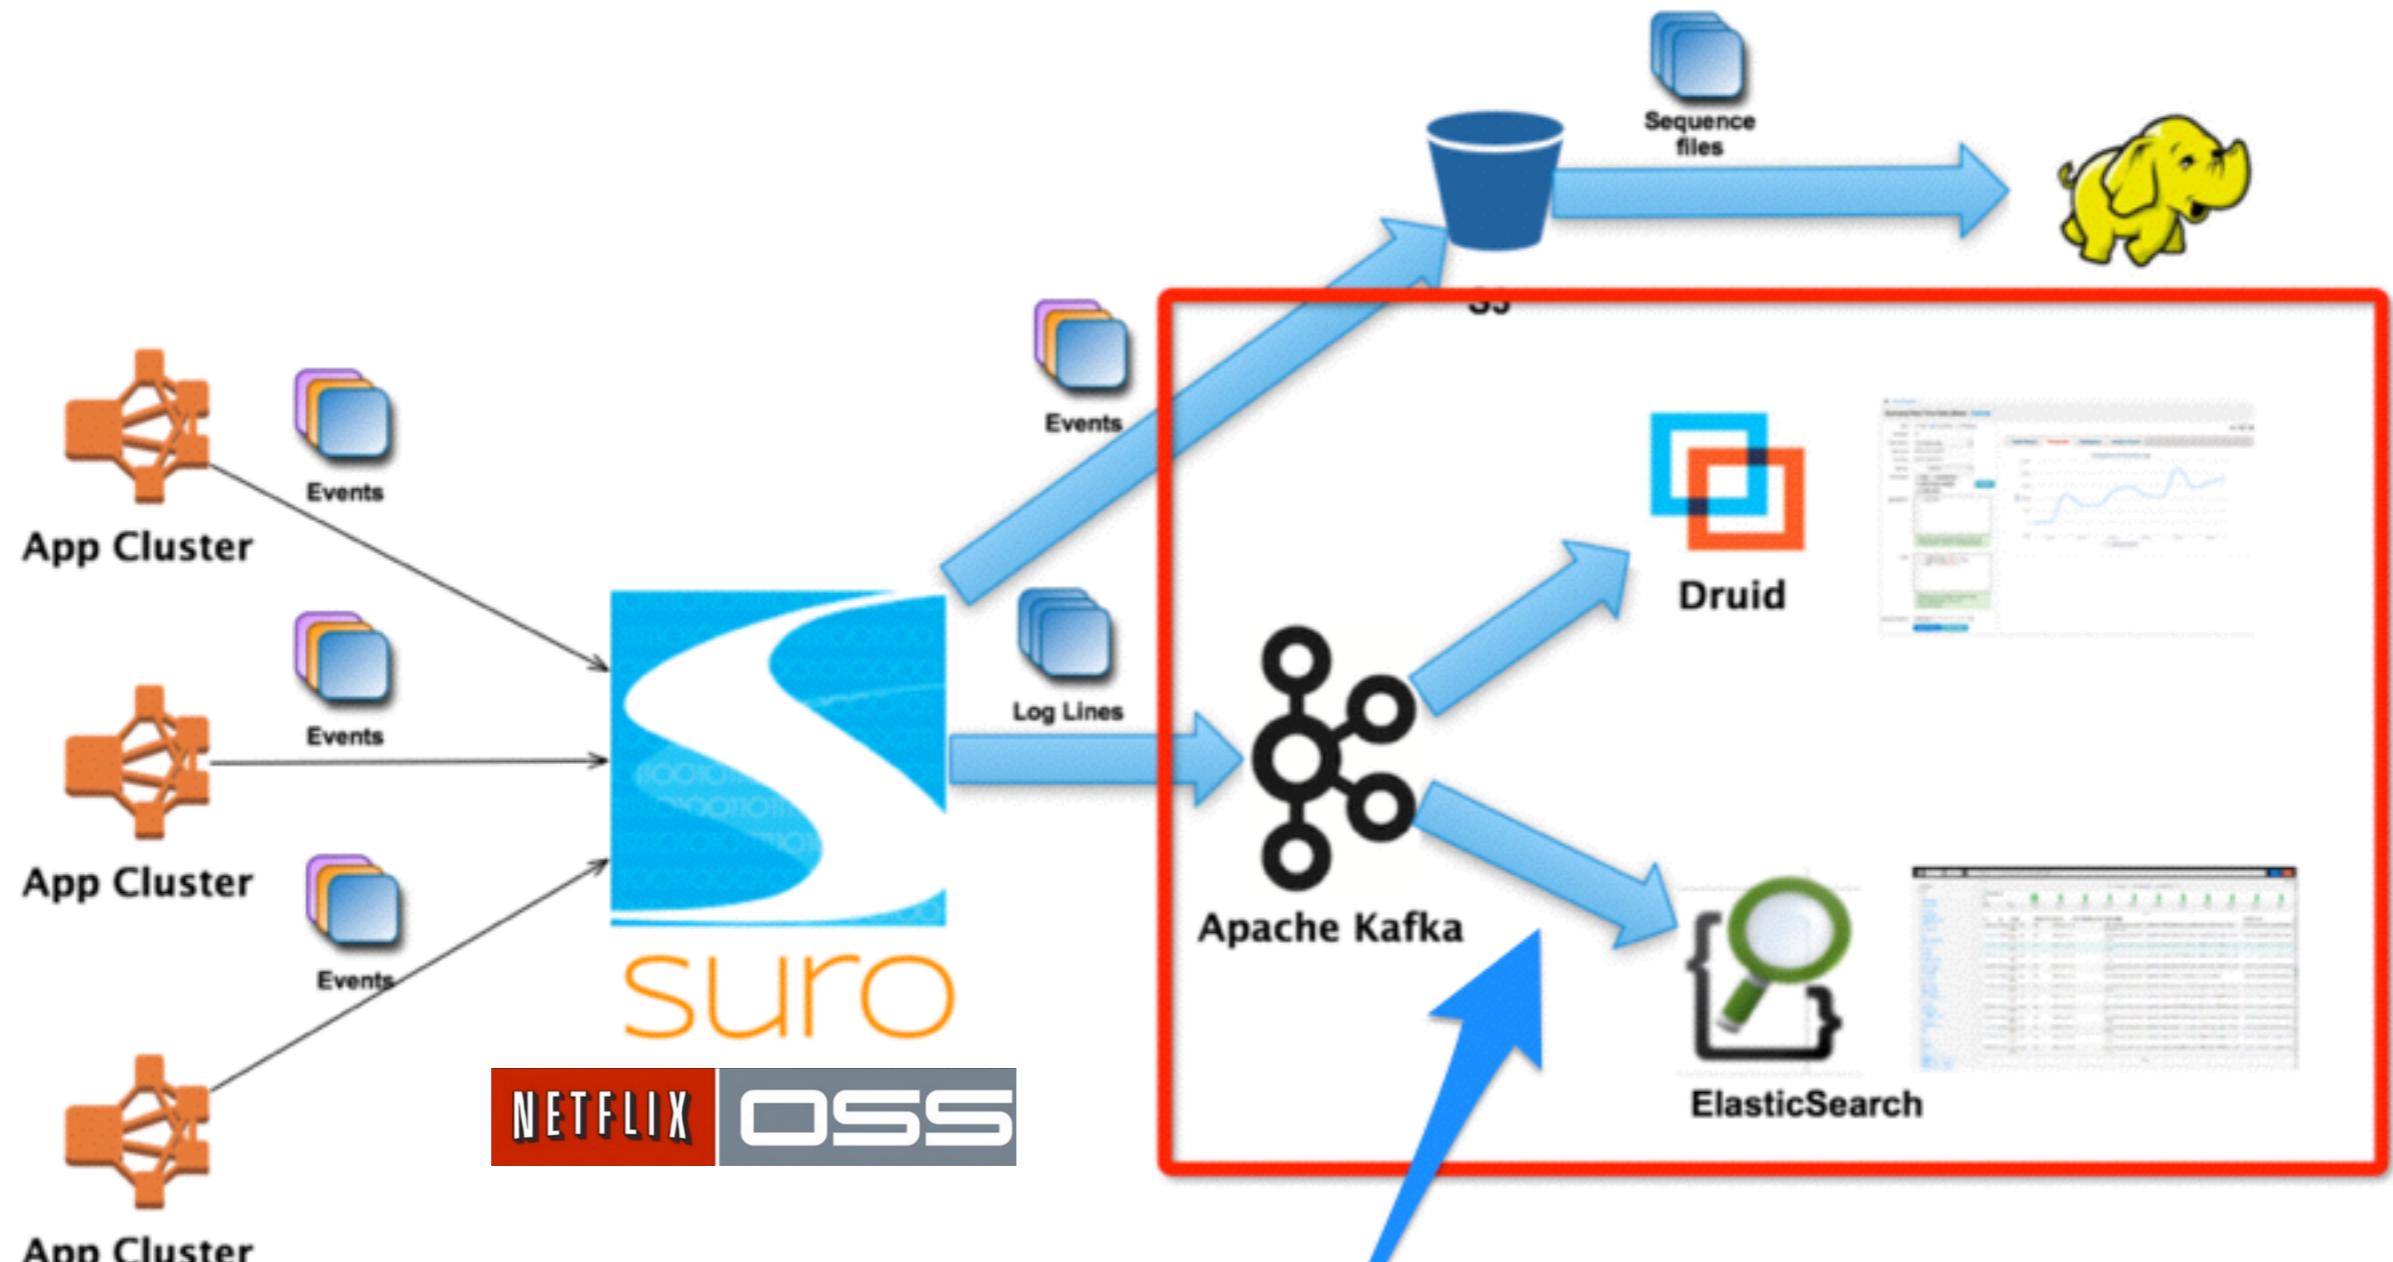
Tolerating outage at different scales



Favor Pulling Over Pushing



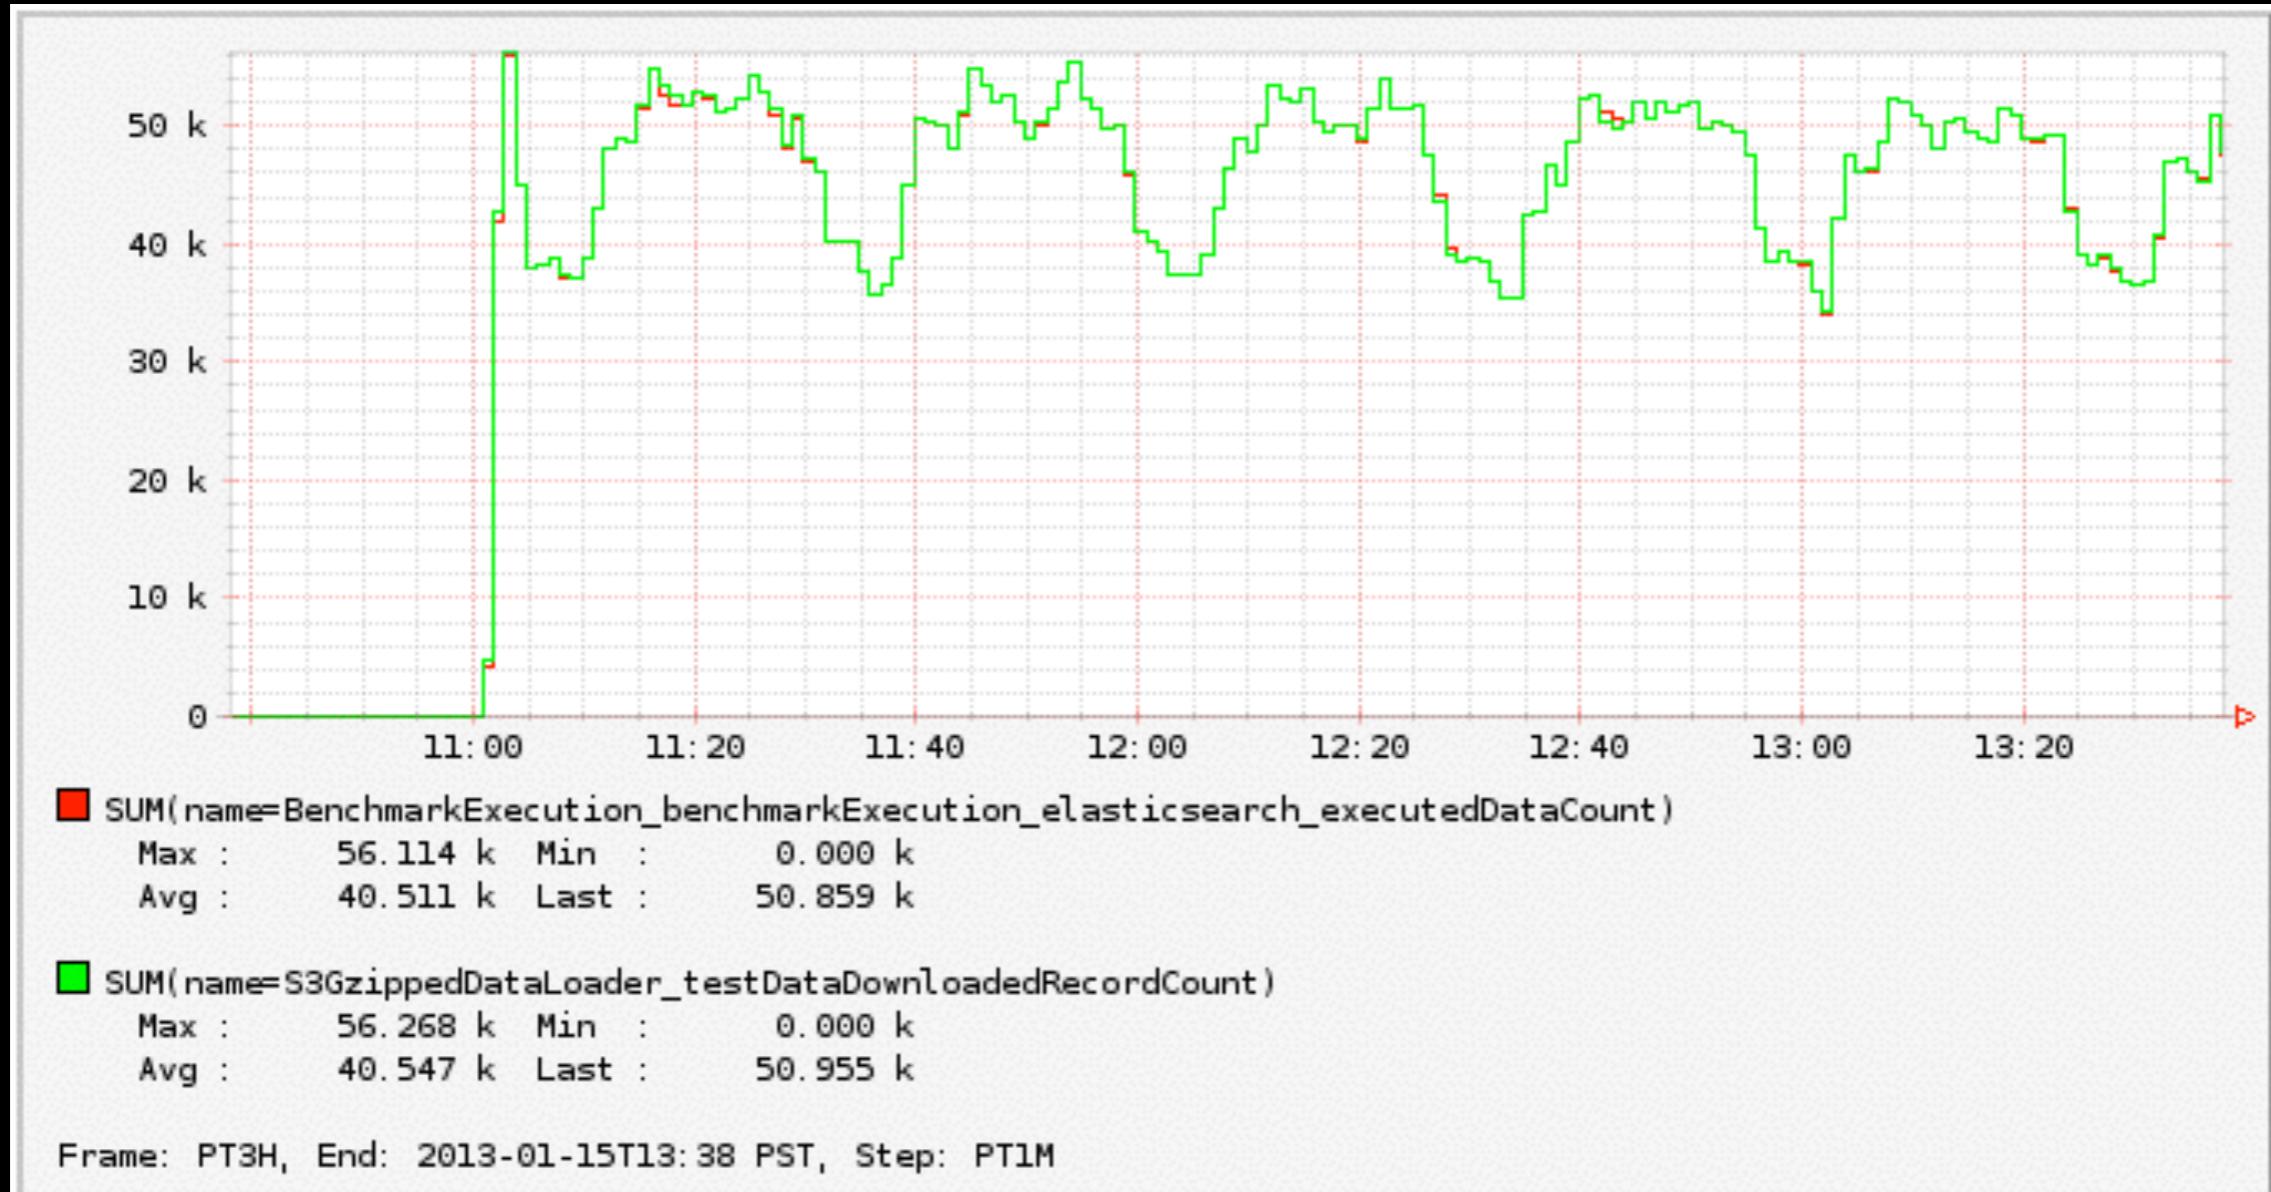




Elasticsearch Kafka River

Choose Config with Data





NETFLIX®

Integrating ES



AMI for Deployment by Asgard



ASGARD prod us-east-1

Home App AMI Cluster ELB EC2 SDB SNS SQS RDS Task

Images in us-east-1 (Virginia)

ID	Name	Description	State	Owner	Creator	Creation Time	Last Used	Package Version
ami-2de3d344	nfelasticsearch-0.90.10-h9.e4350d2-x86_64-201401142121-centos-pv-ebs	name=nfelasticsearch, arch=x86_64, ancestor_name=centosbase-x86_64-201308302051-ebs, ancestor_id=ami-4d185424, ancestor_version=nflx-base-1.2.1-1942837.h806	available	test	jbae	2014-01-14 21:28:05 UTC	2014-01-21 12:01:20 PST	nfelasticsearch-0.90.10-h9.e4350d2/WE-APP-ElasticSearch-Kafka_0_7/9
ami-0720116e	nfelasticsearch-0.90.10-h4.e632b09-x86_64-201401110827-centos-pv-ebs	name=nfelasticsearch, arch=x86_64, ancestor_name=centosbase-x86_64-201308302051-ebs, ancestor_id=ami-4d185424, ancestor_version=nflx-base-1.2.1-1942837.h806	available	test	gmuthusamy	2014-01-11 08:33:32 UTC	2014-01-21 12:01:20 PST	nfelasticsearch-0.90.10-h4.e632b09/WE-APP-ElasticSearch-V0_90_10/4
ami-0330066a	nfelasticsearch-0.90.2-h3.34288b6-x86_64-201401081823-centos-pv-ebs	name=nfelasticsearch, arch=x86_64, ancestor_name=centosbase-x86_64-201308302051-ebs, ancestor_id=ami-4d185424, ancestor_version=nflx-base-1.2.1-1942837.h806	available	test	jbae	2014-01-08 18:27:17 UTC	2014-01-21 12:01:20 PST	nfelasticsearch-0.90.2-h3.34288b6/WE-APP-ElasticSearch-Cass/3
ami-31e3d458	nfelasticsearch-0.90.2-h2.ccbad33-x86_64-201401070019-centos-pv-ebs	name=nfelasticsearch, arch=x86_64, ancestor_name=centosbase-x86_64-201308302051-ebs, ancestor_id=ami-4d185424, ancestor_version=nflx-base-1.2.1-1942837.h806	available	test	jbae	2014-01-07 00:23:09 UTC	2014-01-16 00:01:22 PST	nfelasticsearch-0.90.2-h2.ccbad33/WE-APP-ElasticSearch-Cass/2



Archaius for Configuration



Persisted Properties Console : TEST us-east-1 login

[Home](#) / [FastProperty Explorer](#) / [All Apps](#) / [elasticsearch](#) / [elasticsearch.config](#)

Fast Property Details

[Edit Fast Property](#) [Delete Fast Property](#)

Name:	elasticsearch.config
Value:	<pre>cluster.name: elasticsearch_nccplog_ssd action.auto_create_index: false indices.memory.index_buffer_size: 50% index.number_of_shards: 18 index.number_of_replicas: 2 index.refresh_interval: -1 index.merge.scheduler.max_thread_count: 1 index.merge.policy.max_merged_segment: 1gb index.search.slowlog.level: WARN index.search.slowlog.threshold.query.warn: 10s index.search.slowlog.threshold.fetch.warn: 1s index.analysis.analyzer.default.type: keyword index.store.compress.stored: true index.store.compress.tv: true path.data: /mnt/data/es path.logs: /logs/es bootstrap.mlockall: true transport.tcp.port: 7102 transport.tcp.connect_timeout: 10s</pre>



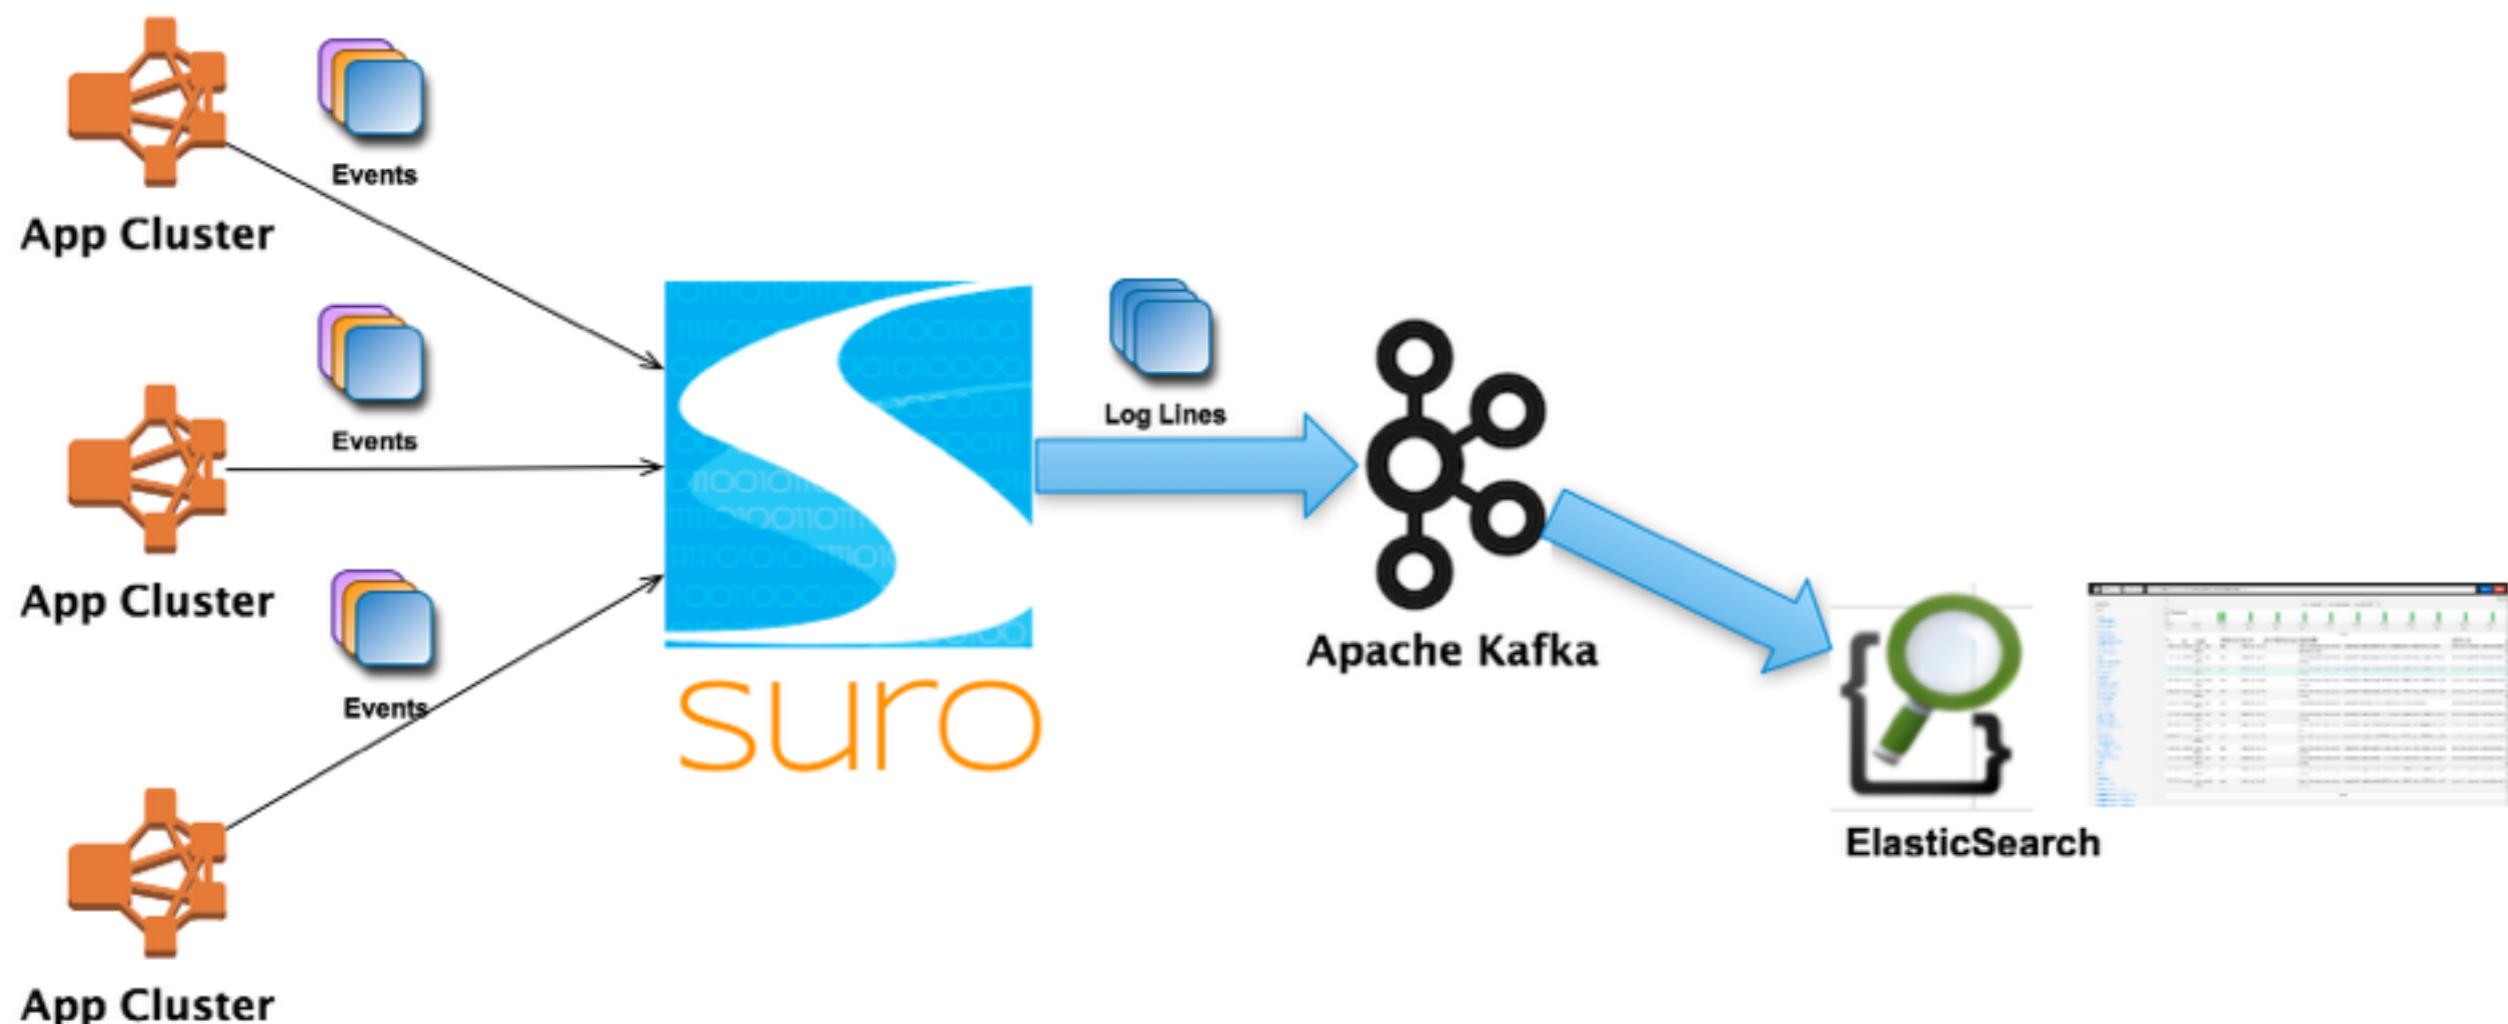
Eureka for Server Discovery



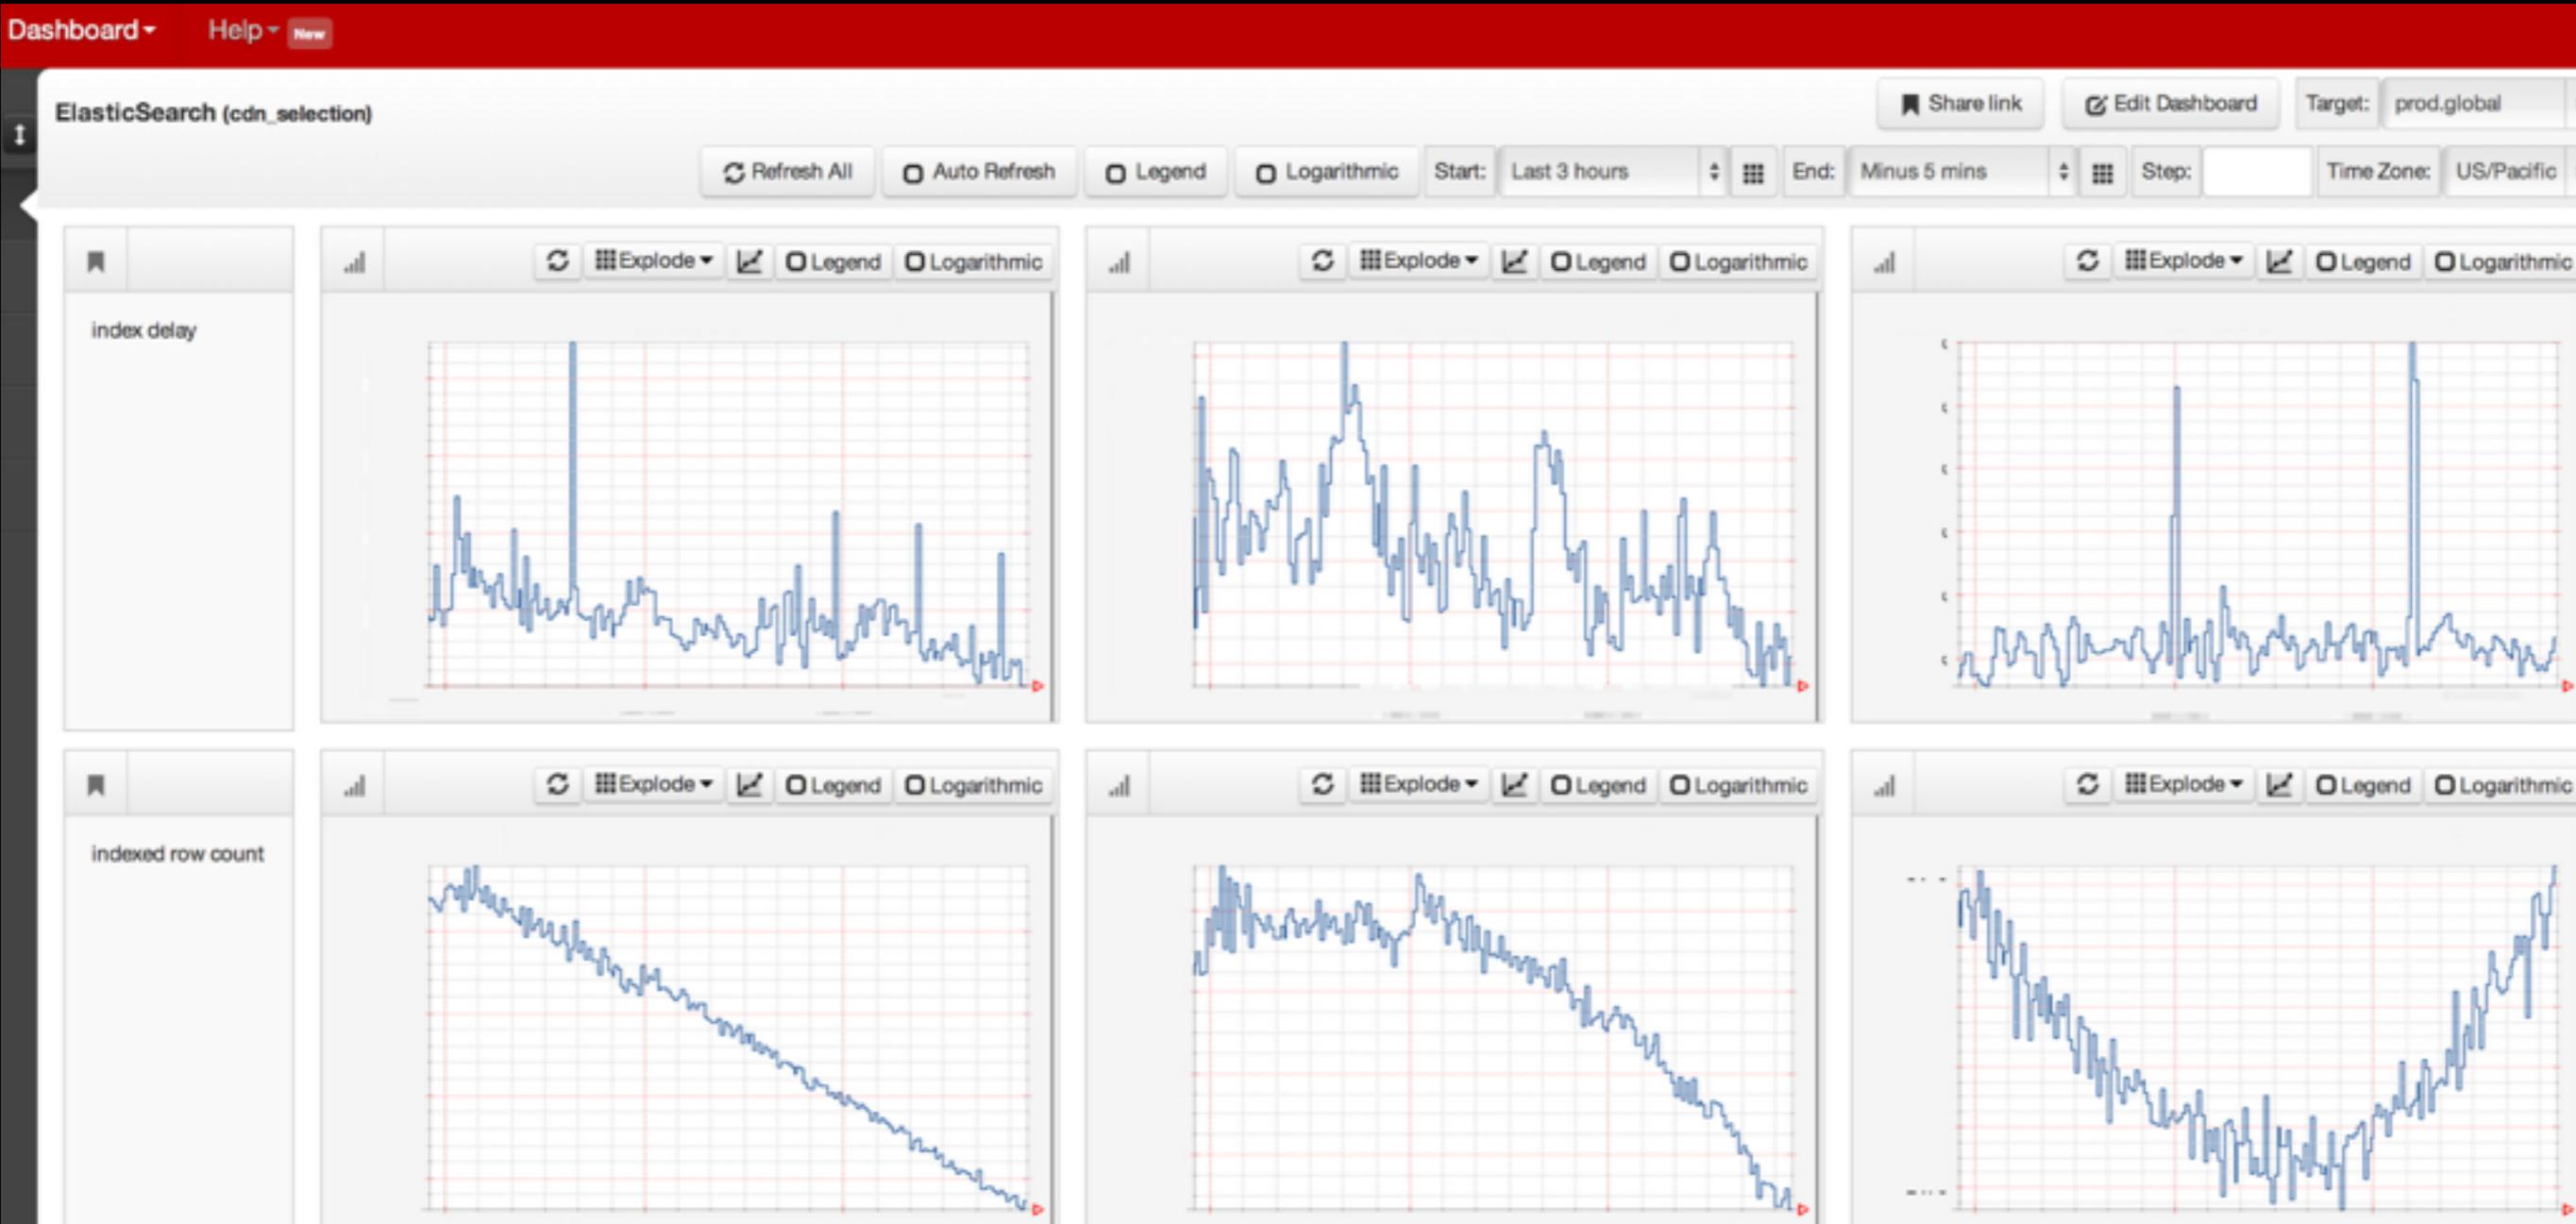
Instances

Total Instances: 3	Run All Health Checks	Terminate Instances								
Load Balancing:	Deregister Instances from ELBs	Register Instances with group's ELBs								
Eureka:	Deactivate in Eureka	Activate in Eureka								
Instance	Zone	State	Launch Time	Package	Ver	Commit	Build	ELBs	Eureka	Health
<input type="checkbox"/>	i-e2174a9e	us-east-1c	InService	2013-12-19 13:04:59 PST	nfelasticsearch	0.90.2	40863c2	8	UP	
<input type="checkbox"/>	i-dad344a5	us-east-1e	InService	2013-12-19 13:45:38 PST	nfelasticsearch	0.90.2	40863c2	8	UP	
<input type="checkbox"/>	i-c77d0fbf	us-east-1d	InService	2013-12-19 13:47:39 PST	nfelasticsearch	0.90.2	40863c2	8	UP	

Suro for Data Delivery



Servo for Monitoring Metrics



Zone-aware Replication



Multi-region Deployment



Multi-region Deployment

Discovery over Cassandra

Region-aware replication



Favor Index Rolling Over TTL



Favor Index Rolling Over TTL

A dedicated service manages index rolling

Uses index template and routing



Worth Trying G1



Worth Trying G1

Not recommended by ES team, but



Worth Trying G1

Not recommended by ES team, but

Has fewer and shorter GC pauses



Worth Trying G1

Not recommended by ES team, but

Has fewer and shorter GC pauses

Occasional SIGSEGV, but it's okay

Simple Majority for Master Election



Simple Majority for Master Election

Split-brain problem



Simple Majority for Master Election

Split-brain problem

`discovery.zen.minimum_master_nodes`

Simple Majority for Master Election

Split-brain problem

`discovery.zen.minimum_master_nodes`

Dynamically updated



Future Work



Future Work

Automatic incremental backup and restore



Future Work

Automatic incremental backup and restore

Auto scaling



Future Work

Automatic incremental backup and restore

Auto scaling

Fully automated deployment



Future Work

Automatic incremental backup and restore

Auto scaling

Fully automated deployment

Support more use cases



We're Hiring



Thank You!

